

eosc-basic-compliance — analysis workflow

What the test suite does, step by step, for every checklist requirement

Mario Reale

26 September 2026, documentation commit 3769a3e

Contents

1. The three stages
2. Stage 1 — collection
3. What the evidence contains — and what it does not
4. Stage 2 — the assessment machinery
5. Per-point decision procedures — points 1, 1R, 2, 3, 4, 5a, 5b, 5c, 6, 7
6. Where this workflow is known to be wrong
7. Reading the output
8. Worked examples — adding a node, changing a URL, a new checklist revision

What the test suite actually does, step by step, when it assesses a node against the **Node Landing Page Verification Checklist v3.0**.

This document is written against the code in [src/basic_check/](#) as it stands, not against the design intent. Where the implementation is weaker than the checklist wording, that is stated rather than smoothed over — the point of this document is that a reviewer can predict the verdict before running the tool, and can tell which verdicts are worth trusting. It was last checked against commit [6423432](#) on 26 September 2026: every [file.py:line](#) reference below points at the current source, and the figures quoted from runs say which run they come from.

Reading order: sections 1–4 describe the machinery shared by every point; section 5 gives the decision procedure for each of the ten points individually; section 6 lists the known ways the machinery can be wrong; section 7 describes the output; section 8 gives three worked examples — adding a node, changing a node's URL, and moving to a new checklist revision — with every command.

1. The three stages

The tool is deliberately split so that **no judgement is ever made against a live website**. Evidence is captured once, written to disk, and every verdict is then derived from that file.

Stage	Command	Reads	Writes	Network
1. Collect	<code>collect</code>	<code>nodes.yaml</code>	<code>results/evidence/<node>.json</code> (personal data masked), <code>results/evidence/screenshots/<node>.png</code>	yes
2. Assess	<code>assess</code>	evidence JSON + <code>checklist/v3.0.yaml</code>	<code>results/results.json</code>	no
3. Report	(part of <code>assess</code>)	the run dict in memory, masked again before rendering	<code>index.html</code> , <code>results.md</code> , <code>results.csv</code> , <code>checklist-v3.0.html</code>	no

`run` is stages 1–3 in one invocation. `points` and `show` are inspection commands and `assess` nothing.


Two consequences of the split matter in practice:

- **A verdict can be re-derived without touching the node's server.** After a checklist reinterpretation or a bug fix, re-running `assess` against the stored evidence regenerates every verdict offline, and the new report can be diffed against the old one.
- **The checks cannot see anything the collector did not record.** If a fact is not in `PageEvidence`, no check can use it — see section 3.

The phases in time, their options, and what to re-run

The three phases, in execution order

Each phase reads what the previous one wrote to disk. Only phase 1 contacts the nodes.



run --update-results-for-node NODE runs the three phases for one node only; every other row is copied unchanged from the stored results.json.

Figure 1. The three phases in execution order: collect contacts the nodes, assess and report work offline, and run performs all three.

The phases run strictly in this order, and each one reads only what the one before it wrote to disk. **Collect** is the only phase that contacts the nodes. It reads `nodes.yaml` and writes `evidence/<node>.json` and a screenshot per node. **Assess** reads that evidence, the checklist and the approved-name list, and sends no request. **Report** is the last step of **assess**: it writes `results.json`, `results.md`, `index.html`, `results.csv` and `checklist-v3.0.html`. **run** performs the three in one invocation. With `--update-results-for-node`, **run** does the same for one node only, and `results.json` supplies every other row unchanged. `points` and `show` only read, so they never need re-running.

Options of each phase. ✓ means the command accepts the option. Every command also takes `--help`.

Option	collect	assess	run	What it does
<code>--nodes, -n <file></code>	✓	✓	✓	Node list to use instead of nodes.yaml
<code>--results <dir></code>	✓	✓	✓	Folder to write to or read from instead of results/
<code>--only <ids></code>	✓	✓	✓	Fetch only these node ids
<code>--skip <ids or names></code>	✓	✓	✓	Leave these nodes out entirely; the report says so
<code>--url <url></code>	✓	✓	✓	Check a page not in nodes.yaml (writes to results/one-off/)
<code>--eosc-page <url></code>	✓	✓	✓	With one <code>--url</code> : that node's own eosc.eu page, for point 4
<code>--node <id or name></code>	✓	✓	✓	One configured node at the alternative page given with <code>--url</code>
<code>--delay <s></code>	✓	—	✓	Seconds between hosts (default 2.0)
<code>--depth 0/1/2</code>	✓	—	✓	How far to follow links (default 1)
<code>--max-children <n></code>	✓	—	—	Linked pages followed per node at depth 1 (default 8; <code>run</code> always uses 8)
<code>--fetch-budget <n></code>	✓	—	✓	Depth 2 only: ceiling on second-hop requests (default 60)
<code>--checklist, -c <file></code>	—	✓	✓	Checklist revision to apply (default checklist/v3.0.yaml)
<code>--approved-names <file></code>	—	✓	✓	Approved-name list for point 3 (default checklist/approved-names.txt)
<code>--no-approved-names</code>	—	✓	✓	Use no name list at all
<code>--strict-separators</code>	—	✓	✓	Match separators in approved names literally
<code>--run <label></code>	—	✓	✓	Label recorded in every report (default: UTC time)
<code>--update-results-for-node <node></code>	—	✓	✓	Update only that node's row in the stored results; <code>assess</code> re-judges its saved evidence, <code>run</code> fetches it again
<code>--accept-error</code>	—	—	✓	With <code>--update-results-for-node</code> : merge the row even if the fetch failed

What to launch again after a change. Which phase has to be repeated depends on which input changed. A change to [nodes.yaml](#) needs a new collection for the nodes concerned. A change to the checklist or to the name list needs only a new assessment, which is offline. `--update-results-for-node` refuses when the checklist or the approved-name list differs from the stored results', because a row judged by other rules would not be comparable with the others. In those cases, every row is assessed again from its stored evidence.

You change	Launch again	Sites contacted	The other nodes' rows
Add a node: an entry in <code>nodes.yaml</code> and its approved name in <code>checklist/approved-names.txt</code> (and <code>-scoped.txt</code>)	<code>collect --only <new id></code> , then <code>assess</code> without <code>--only</code>	the new node only	Re-judged, offline, from their stored evidence, because the approved-name list used for point 3 changed
Add a node whose approved name is already in <code>checklist/approved-names.txt</code>	<code>run --update-results-for-node <new id></code>	the new node only	Kept unchanged; the new row is inserted in <code>nodes.yaml</code> order
Change a node's landing page URL (<code>basic-check --update-nlp <id> <URL></code> , or by hand)	<code>run --update-results-for-node <id></code>	that node only	Kept unchanged
Go back to the node list as downloaded (<code>basic-check --restore-default-config</code>)	For each node it lists as <code>URL</code> restored: <code>run --update-results-for-node <id></code> . If it lists nodes added back or removed: <code>assess</code>	the restored nodes only	Kept unchanged; re-judged offline if nodes were added back or removed
Change the reference checklist (a new <code>checklist/vX.Y.yaml</code> , made the default in <code>DEFAULT_CHECKLIST</code> or passed with <code>-c</code>)	<code>assess</code>	none	All re-judged, offline, from the stored evidence

`assess` alone is not enough after a URL change. The node's evidence still comes from the old page, so `assess` warns and both reports carry an **Evidence from a different URL** banner. After a checklist change, a new `collect` (or a full `run`) is needed only if the new revision asks for something the collector does not record (section 3), or if you want a fresh capture date. For the full procedure of each case, with every file and command, see the worked examples in section 8. Each command writes to `results/`, the published run, by default. Add `--results /tmp/copy` to work on a copy first. The published run covers all thirteen nodes, so no `--skip` is needed for the same scope. With an explicit `--results`, `--only` also narrows the report to those nodes, which is why the table runs `assess` without it.

2. Stage 1 — collection

Per node, in order:

1. **robots.txt for the landing page host** (`_robots`, `fetch.py:194`). Fetched with a plain HTTP client (`httpx`, not the browser), with the tool's own User-Agent string. Three outcomes:
 - o disallowed → `error = "not fetched: robots.txt disallows it"`, evidence record returned immediately, **every point becomes ERROR**. This is not hypothetical: on a trial run on 25 September 2026, BBMRI-ERIC's new `dev3`. address served `User-agent: * / Disallow: /`, and all ten of its points were `ERROR`;

- o HTTP ≠ 200 → treated as "nothing disallowed", proceed;
 - o unreachable → proceed with one request, and the reason is recorded in `robots_note`. A network failure is not read as permission, but it is also not read as a prohibition.
- 2. **Render in a real browser.** Chromium via Playwright, viewport 1440×1000, locale `en-GB`, `Accept-Language: en-GB,en;q=0.9`. Navigation waits for `domcontentloaded` (45 s cap), then gives `networkidle` a bounded 2.5 s chance, then a fixed 1.2 s settle. Institutional sites run analytics that never let the network go idle, so the wait is bounded and whatever has rendered is taken. Rendering rather than plain HTTP matters: several of these landing pages are client-side rendered, and an HTML fetch would see an empty shell and report the node as having no content — a false accusation produced by the tool.
- 3. **Record status, final URL and redirect chain.** Redirects are captured from response events (first 10 kept).
- 4. **Extract the evidence** from the rendered DOM (`_extract`, `fetch.py:211`). `script`, `style`, `noscript` and `template` are removed first. Then:
 - o `full_text` — all body text, whitespace-collapsed;
 - o `main_text` — text of the first of `main`, `[role=main]`, `article`, `#content`, `.content`, falling back to `body`. Kept separate so a cookie banner or a language switcher cannot sway language detection;
 - o `links` — every `a[href]` except `javascript:` and bare `#`, resolved to absolute URLs; if the anchor has no text, a nested `img`'s `alt` is used as its label;
 - o `images` — every `img` (`src`, `alt`, `aria-label`, `title`, `class`) plus every `svg` (its text content as the label);
 - o `controls` — every `button`, `a[href]`, `input[type=submit]` and `[role=button]` that has a label, taken from its text, its `aria-label` or its `value`;
 - o `title`, `lang_attr` (the `html` element's `lang`), `meta_description`.
- 5. **Screenshot** the viewport (not the full page) to `results/evidence/screenshots/<node>.png`. Nothing in the assess stage reads the screenshot; it exists for the human reviewer. It captures the visible viewport, not the full page, so it cannot be used to confirm the absence of anything — a footer link is real but off-image.
- 6. **Depth 1, if requested.** `select_children` (`fetch.py:414`) picks which links to follow. A link qualifies only if it is `http(s)`, is not a binary by extension, is on the node's own host or a related host, and its combined link text and URL path matches the vocabulary of a point that a second page can settle. The vocabulary is defined once in `patterns.py` and consumed by both the crawler and the checks, so the crawler cannot fail to fetch something a check would have accepted.

The purposes, in the order they consume budget:

Order	Point	Vocabulary
1	5b	Acceptable Use Policy, AUP, terms of use/service, conditions of use, user agreement
2	5c	user access policy, UAP, access policy, conditions of access, access conditions

Order	Point	Vocabulary
3	6	helpdesk, help desk, service desk, support, ticket, contact, get in touch
4	2	about, who we are, our mission, mission

Caps: `MAX_PER_PURPOSE = 2`, `MAX_CHILDREN = 8` per node (overridable with `--max-children`), 1.2 s between child requests. Each child page carries `selected_for`, the list of points it was fetched for, and every link that was matched but dropped is recorded in `children_skipped` with the reason. That list is the difference between "checked and absent" and "never looked".

Child pages are fetched more cheaply than the landing page: 25 s timeout, 600 ms settle, no `networkidle` wait, no screenshot.

7. **Depth 2, if requested.** Only from children that rendered — a child that failed has no trustworthy links. Meaner caps: `MAX_PER_PURPOSE_D2 = 1`, at most 2 grandchildren per child, a run-wide fetch budget of 60, 1.5 s between requests. A URL already fetched for this node is never fetched twice. `robots.txt` is consulted per host at this level too, cached one fetch per host per run, because following links reaches hosts the landing page never named.

When depth 2 was used, the assess stage runs **both** the full assessment and an assessment of the same evidence with the second hop stripped (`without_depth_2`), storing the latter as `results_depth_1`. The report can then show what the extra requests actually changed instead of asserting they were worth making.

8. **Write** the whole record as JSON.

3. What the evidence contains — and what it does not

PageEvidence (fetch.py:126) holds: `node_id`, `node_name`, `requested_url`, `fetch_at`, `final_url`, `http_status`, `redirect_chain`, `robots_allowed`, `robots_note`, `title`, `lang_attr`, `main_text`, `full_text`, `links`, `images`, `controls`, `meta_description`, `error`, `screenshot`, `crawl_depth`, `children`, `children_skipped`, `crawl_note`.

Image holds only `src`, `alt`, `aria_label`, `title`, `css_class`, `inline_svg`. Link holds `href`, `text`, `visible`. Control holds `text`, `href`, `tag`.

No geometry, visibility or computed style is captured anywhere. `getBoundingClientRect`, `offsetWidth` and `getComputedStyle` do not appear in `src/`. This is the single most consequential gap in the design, and it is why point 3 can never be closed automatically: the checklist asks whether a logo is *visibly displayed*, and the tool has recorded nothing about visibility. It knows an element exists in the markup. It does not know whether a human would see it.

Likewise, a logo delivered as a CSS `background-image`, as an SVG sprite reference, or from a file whose name does not contain `eosc`, is invisible to the tool. Its absence from the evidence is not evidence of absence.

Every string in the evidence is saved with personal data masked (`privacy.py`, through `write_evidence` in `fetch.py`). Personal addresses become `XXXXXX@domain`, phone numbers become `+CC XXXXXX`, and a name next to a masked address becomes `XXXXXX`. Role mailboxes are kept. The checks therefore run on masked text, which costs nothing: point 6 needs a helpdesk, and no check uses a person's address or number. Re-assessing the 24 September evidence after masking gave the same verdict in all 120 cells, and that run was

masked this way at commit [ada1b4a](#). Run [web-4](#), which replaced it as the published run on 26 September, was masked as it was collected. The reports are masked a second time as they are written ([write_all](#) in [report.py](#)), so a report rebuilt from evidence captured before masking existed is masked too. Screenshots are not masked.

4. Stage 2 — the assessment machinery

`checks.run_all(ev, approved_names, expected_eosc_page)` calls exactly ten functions in checklist order and returns exactly ten `Result` records. The mapping between checklist point ids and check functions is asserted to be total in both directions by the test suite, so a point cannot be silently unassessed and a function cannot assess a point that no longer exists.

A `Result` carries `point_id`, `title`, `verdict`, `message`, `evidence` (a list of strings), and `reviewer_action`.

The verdict vocabulary

Verdict	Meaning
PASS	Satisfied, and the tool can show why
FAIL	Violated, and the tool can show why
MANUAL_REVIEW	A human must decide; evidence is attached to make that fast
ERROR	The tool could not assess at all (page unreachable, robots disallowed)

`MANUAL_REVIEW` is the load-bearing one. Several checklist points turn on phrases like "clearly state", or quantify over "all research resources offered by the Node". Neither is settleable by inspecting one page, and emitting `PASS` or `FAIL` on them would be a guess dressed as a verdict. A wrong `FAIL` against a named organisation is expensive to retract, so where the tool cannot know, it says so and hands over the evidence.

The first step of every check

Every one of the ten functions begins the same way:

```
if ev.error: -> ERROR
```

So a node whose landing page could not be fetched at all returns ten `ERROR` cells, never a mix of passes and a failure.

Shared helpers

Helper	Line	What it does
<code>_match(text, patterns)</code>	58	Returns the matched substring for each pattern that hits
<code>_link_hits(ev, patterns)</code>	67	Links whose " <code>{text} {href}</code> " matches any pattern
<code>_fmt(link)</code>	76	Renders a link as " <code>label</code> " -> <code>href</code> for the evidence list

Helper	Line	What it does
<code>_is_host(netloc, domain)</code>	164	Exact host or true subdomain . Rejects <code>myeosc.eu</code> and <code>not-eosc.eu</code> , which a naive <code>endswith</code> would accept
<code>_verify_child(ev, point, ...)</code>	571	The child page that was followed for this point, or <code>None</code>
<code>_policy_check(...)</code>	578	The whole decision procedure shared by points 5b and 5c
<code>render_warning(ev)</code>	—	Reason to distrust <i>absence of text</i> on this page
<code>link_collection_warning(ev)</code>	—	Reason to distrust <i>absence of a link</i>

The two absence gates — and why there are two

Concluding "absent" from an incomplete capture is the most damaging mistake this kind of tool can make, because the output is an accusation against a named organisation. So absence-based `FAILs` are gated. But the two kinds of absence need different gates.

`render_warning` asks *did the page render enough to read?* It fires when a cookie-consent overlay dominates the captured text (consent markers present and `main_text` under 1200 chars), when fewer than 25 links were captured, or when `main_text` is under 600 chars. This is the right question for a judgement about prose and the **wrong** question for a link.

`link_collection_warning` asks only *did the DOM arrive?* It fires when no links at all were captured, when $\text{len(links)} + \text{len(images)} + \text{len(controls)} < 10$ (`DOM_ELEMENT_FLOOR`), or when `full_text` is under 500 chars (`EMPTY_DOCUMENT_CHARS`). Text length is not evidence about DOM delivery — a small page is not a truncated one.

The distinction came from a real case. One node in this set serves 530 characters of main text behind a cookie banner but 20 anchors — exactly the 20 in the server's HTML, and 21 in the rendered DOM both before and after accepting consent. Gating a link-absence `FAIL` on text length let a genuinely verified absence escape as "cannot conclude". The absence of an `eosc.eu` link there is a fact about the node, not an artefact of collection.

Where a `FAIL` stands despite a render warning, `_render_note` appends the caveat to the evidence: the verdict is not suppressed, but the reviewer is told a consent banner was in the way so the finding can be spot-checked.

5. Per-point decision procedures

Ten points, ten functions. For each: the steps taken, the branch conditions in evaluation order, and the verdict each branch yields.

Point 1 — NLP publicly accessible, or reachable via EOSC AAI

`check_1`. Decided almost entirely on HTTP status.

Steps: read `http_status`; measure `len(full_text)`; on a 401/403, search `full_text` for login vocabulary and `final_url` for AAI markers, then search `full_text + title` for bot-wall wording.

#	Condition	Verdict
1	error set	ERROR
2	2xx and <code>full_text</code> > 200 chars	PASS — branch (a) satisfied
3	401/403 and a login affordance found	MANUAL_REVIEW — branch (b) may apply, but whether that login is EOSC AAI cannot be read off the page
4	401/403, no login affordance	MANUAL_REVIEW
5	404 or 410	FAIL — the registered URL does not resolve to a page
6	≥ 500	MANUAL_REVIEW — often transient, not recorded as a failure without a retry
7	any other ≥ 400	MANUAL_REVIEW
8	2xx but ≤ 200 chars	MANUAL_REVIEW — may need JavaScript the tool did not run, or be a shell

Branch 4 deserves its own note, because it looks like it should be a failure and is not. A bare 403 to an automated client is far more often bot mitigation than an access policy — genuine access control almost always redirects to a login. This was not hypothetical: a node in this set served HTTP 200 to the tool and then 403 once the crawl made a handful more requests. Reporting that as "not publicly accessible" would have been the tool blaming a node for its own request rate. The same node returned 200 again hours later to a run from an address that had made no requests that day, settling all four of its review cells — so the abstention was not merely cautious, it was correct, and a FAIL would have stood as a wrong verdict about a named organisation. The evidence records any bot-protection wording found (`cloudflare`, `just a moment`, `ray id`, `captcha`, ...) and the reviewer action is to open the URL in an ordinary browser.

Point 1R — resources behind the landing page are public or behind EOSC AAI

`check_1R`. Returns **MANUAL_REVIEW** unconditionally (unless `error`).

Steps:

1. Take the landing page host from `final_url`.
2. Collect every link whose host differs from it and is not `eosc.eu` or a subdomain (via `_is_host`).
3. Group those by host, listing up to 12 hosts with up to 3 distinct link labels each, so the reviewer sees distinct destinations rather than 90 URLs.
4. Search `full_text` and all `hrefs` for AAI markers (`eosc-aai`, `aai.eosc`, `myaccessid`, `egi check-in`, `aai.egi.eu`, ...) and list any found.
5. If depth ≥ 1, report how many followed child pages were served anonymously, and name up to 5 that were not.

The point quantifies over every resource reachable through the landing page, including through intermediate pages, and whether a given login is genuinely EOSC AAI compliant is settled during EEN enrolment rather than by reading HTML. One level of crawling narrows this; it cannot close it. The output is a worklist, not a verdict.

Point 2 — scope, intended users and responsible organisation stated

`check_2`. Returns **MANUAL_REVIEW** unconditionally (unless `error`).

Steps:

1. Take the first 600 characters of `main_text`, falling back to `full_text`.
2. Add the `meta` description if present (first 200 chars).
3. Find organisation-like names by regex: up to four capitalised words followed by a legal-entity or institutional suffix — `ERIC`, `e.V.`, `GmbH`, `Ltd`, `Foundation`, `Institute`, `Institut`, `University`, `Universit[éeà]`, `Consortium`, `Association`, `Council`, `Agency`, `Centre`, `Center`, `CNRS`, `CNR`, `CSC`. Up to six are listed; if none match, that is stated explicitly.
4. If an About page was followed at depth 1 and rendered, quote its opening 260 characters — **clearly marked as one level down**, because the checklist asks the *landing page* to state these things and an About page cannot satisfy it on the page's behalf.

Whether the prose conveys scope, intended users and the responsible organisation to a researcher is a reading judgement. The check's job is to make that reading take thirty seconds instead of five minutes.

Point 3 — EOSC logo and official node name visible

`check_3`. Returns **MANUAL_REVIEW** on both branches (unless `error`).

Steps:

1. For every captured image and inline SVG, build the haystack "`{src} {alt} {aria_label} {title} {css_class}`" — but with the **host part of `src` removed**, so a node served from a host containing "eosc" does not match on every locally-hosted image. The path and query survive, so `eosc-node-final.webp` still counts. "eosc" must then appear as a token, not inside a longer word: `geoscience` and `neoscope` do not match, while the official lockups `EOSCNode_Finland.jpg` and `EOSCNodeBBMRIERIC_ColourPos.png` do, because CamelCase is treated as a word break. Separately, an asset genuinely served from `eosc.eu` counts via a host check, so a cross-host node logo — correct behaviour — is still found, while `myeosc.eu` is not.
2. For each hit, record the kind (`img` or `inline SVG`) and a label, preferring `alt`, then `aria_label`, then `title`, then the filename.
3. Take the names that apply to *this* node — those written as `node-id: Name`, then any unscoped ones — and test each against `full_text` in file order. The list is `checklist/approved-names.txt` unless `--approved-names` replaces it or `--no-approved-names` suppresses it. Each word of a name is matched literally and case-insensitively, and never inside a longer word or hyphenated token; the separator *between* words is flexible, so whitespace and the glyphs `|`, `-`, `—`, `;`, `/`, `·` are interchangeable — unless `--strict-separators` is passed, which restores literal glyph matching (and is recorded in the output, so a reader knows which rule was in force). Record the first match, what the page actually wrote, and whether the name was scoped to this node. If the body is empty (a 403, say), record that the name

was not looked for rather than that it is absent. If nothing matched, also record how many times the longest leading phrase shared by all the candidate names — [EOSC Node](#), for the official list — occurs in the body, so the reviewer knows whether there is anything to look at.

#	Condition	Verdict
1	At least one EOSC-referencing image asset	MANUAL REVIEW — the logo requirement is <i>likely</i> met
2	None	MANUAL REVIEW — explicitly not proof of absence

Neither branch can be closed, for two separate reasons that are worth keeping apart:

- **A missing capability.** "Clearly and visibly shown" is a question about rendered geometry, and section 3 explains that no geometry was captured. This is not a judgement call the tool declines to make; it is data the tool does not have.
- **A name the tool can check, against a list it cannot validate.** The official Tripartite list is now committed at [checklist/approved-names.txt](#) and used by default, so the name half is assessed — but the tool cannot tell whether that file is current, and a non-match may mean the page is wrong, the name is in an image, or the list is stale. [results.json](#) records which list was in force under [approved_names](#), including whether it was the committed default, and both reports state it in prose, so a [REVIEW](#) here cannot be misread as a checked-and-clean name. [--no-approved-names](#) skips the half entirely, and the reports say that too.
- **A weaker claim from an unscoped list.** A name given without a [node-id:](#) prefix applies to every node, so a match shows the string is on the page but not that it is that node's own name. The evidence line says which of the two claims it is making. Scoping each name to its node is what makes the stronger claim available, and [checklist/approved-names-scoped.txt](#) does exactly that. It is deliberately not the default: the names in it are official, but the mapping from each name to a node id in [nodes.yaml](#) was derived here and is not part of the Tripartite file, so that mapping is what a reviewer should challenge first.

Branch 1 used to over-match badly; the history is in section 6.

Point 4 — link to the node's own page on [eosc.eu](#)

[check_4](#). The point that produces most of the [FAILs](#) in practice: all 6 in the published run of 24 September 2026, and 6 of the 7 in the run of 21 September, the seventh being point 6 for a node with no contact route of any kind on its landing page.

Steps: for every link, parse the URL; keep only those whose host is [eosc.eu](#) or a true subdomain ([_is_host](#), so [myeosc.eu](#) does not qualify); strip the trailing slash from the path; sort into three buckets —

- **node_pages** — path contains [building-the-eosc-federation](#) and has a non-empty tail after it;
- **index_only** — path contains [building-the-eosc-federation](#) with nothing after it;
- **other_eosc** — any other [eosc.eu](#) link.

#	Condition	Verdict
1	node_pages non-empty	PASS (up to 3 links quoted)
2	index_only non-empty	FAIL — the index is explicitly excluded by the checklist

#	Condition	Verdict
3	<code>other_eosc</code> non-empty	FAIL — links to <code>eosc.eu</code> but not to the dedicated page; the homepage is excluded
4	none, and <code>link_collection_warning</code> fires	MANUAL_REVIEW — too little captured to conclude absence
5	none, DOM complete	FAIL — no link to <code>eosc.eu</code> found, with any render caveat appended

When `nodes.yaml` records the node's `eosc_page`, that URL is passed in as `expected_page` and appended to the evidence of branches 2, 3, 4 and 5 as "the node's dedicated page exists at ... but is not linked from here". Naming the exact page turns "something is absent" into a one-line fix the node operator can action.

What a **PASS** here does not establish. Three things, all of which a reviewer should know before quoting one:

- **The target is never requested.** The check matches the URL's *shape* — host, path prefix, non-empty tail. A link to a node page that has since been deleted would **PASS** identically. Confirming the destination is a manual step.
 - **Placement is not recorded.** A link in the footer and a link in the body are indistinguishable in the evidence. The checklist requires a link and says nothing about prominence, so footer placement is compliant — but the evidence cannot tell you which you are looking at. EUDAT's link, for instance, is the fourth item in a "Navigation & Legal" footer column, 89% of the way down the page.
 - **The screenshot cannot corroborate it.** Captures are viewport-only (`full_page=False` in `fetch.py`), so anything below the fold — which includes every footer link — is absent from the saved image. A reviewer checking a **PASS** against the screenshot may conclude the tool invented the link. This has happened.
-

Point 5a — purpose description for research resources

`check_5a`. Returns **MANUAL_REVIEW** unconditionally (unless `error`).

The evidence is two numbers: the count of outbound links, and the length of `main_text`. The point quantifies over "all research resources offered by the Node", which cannot be enumerated from the landing page, and the checklist separately allows the description to live in the resource's EOSC Catalogue entry rather than on this page. There is no branch on which the tool could honestly resolve it, so it does not pretend to.

Points 5b and 5c — AUP and UAP accessible

`check_5b` and `check_5c` both delegate to `_policy_check` with different vocabularies. Identical procedure, so it is described once.

Steps:

1. `_link_hits` — links whose label or address matches the policy vocabulary. The address is read twice, raw and with its separators turned into spaces (`patterns.link_haystack`), so `/geant-node-acceptable-use-policy/` matches `acceptable\s+use\s+police`. The crawler selects links with the same function.
2. `_match` on `full_text` — mentions of the vocabulary in the prose.
3. If links were found, look for the child page fetched for this point.

#	Condition	Verdict
1	Link found, no child fetched	PASS , stated as "a pointer, not a verified document", with the reason: depth 0, a PDF or other document, a cap reached, another site, or not selected at collection
2	Link found, child returned 404/410	FAIL — the policy is not accessible
3	Link found, child not ok for another reason	MANUAL_REVIEW — may be a bot restriction rather than a real problem
4	Child fetched, but main_text < 400 chars or no policy wording	MANUAL_REVIEW — reads like a navigation stub, not a policy
5	Child fetched, substantial, policy wording present	PASS
6	No link, but the vocabulary appears in the page text	MANUAL_REVIEW — mentioned but not followable
7	Nothing at all	MANUAL_REVIEW — deliberately not a FAIL

Branch 1 is weaker than it looks, and the message says so: a link labelled "Acceptable Use Policy" pointing at a 404 satisfies "there is a link" while failing the actual requirement, which is that the policy be *accessible*. Running at `--depth 1` converts branch 1 into one of branches 2–5 for ordinary pages on the node's own site. It cannot do so for a PDF or a page on another site, and the reviewer action says so rather than suggesting a re-run that would not help.

Branch 4's "policy wording" test is **POLICY_BODY_PATTERNS**: **must not**, **you may/must/shall/agree**, **permitted**, **prohibit**, **terms**, **policy**, **conditions**, **responsib**, **comply**, **authorised**, **authorized**. It exists to tell a real policy document apart from a navigation page that merely has the word "policy" in its link text.

Branch 7 is not a **FAIL** because the checklist permits the policy to be reached via each resource's entry in the EOSC Catalogue, which this tool does not follow. Failing a node for an absence the checklist does not require to be on this page would be the tool enforcing a stricter rule than the one it is checking.

Point 6 — means of contacting the node helpdesk

check_6. Two tiers: wording that identifies a helpdesk, then general contact vocabulary. The word "support" on its own is only a reason to follow a link. On 24 September 2026 it labelled a funding programme (Czechia), a EuroHPC proposal service (EBRAINS) and a service overview (BBMRI-ERIC), and all three had passed.

Steps:

1. Collect **mailto:** links.
2. **_link_hits** with the full contact vocabulary.
3. From those two sets, keep the ones that **identify a helpdesk**. That means the label or address matches **HELPDESK_STRONG** (helpdesk, service desk, ticket, a support team, request or portal,

user/technical support, contact/ask/get support, a [support@](#) or [support\[at\]](#) address, ServiceNow). A host whose first label is [hd](#), [support](#) or [helpdesk](#) also counts, and so does a mailbox such as [support@](#), [helpdesk@](#), [x-helpdesk@](#) or [x@helpdesk.....](#)

4. Otherwise read **every** page followed for point 6, not only the first. A page counts if it matches [HELPDESK_ON_PAGE](#) (helpdesk, service desk, ticket, a support address, ServiceNow) or gives a helpdesk mailbox. Qualified "support" phrases are strong in a link label, but not in page prose: they describe services there.

#	Condition	Verdict
1	Landing-page link or mailto that identifies a helpdesk	PASS
2	Contact or bare "support" link, a followed page names a helpdesk or gives a helpdesk mailbox	PASS — quoting the wording and the addresses
3	Contact or bare "support" link, followed pages ok , none identifies a helpdesk	MANUAL_REVIEW — says when a "support" label was the only lead, and lists addresses and links not followed
4	Every followed page returned 404/410	FAIL — the route offered does not work
5	Generic contact link, nothing further established	MANUAL_REVIEW
6	No contact route, link_collection_warning fires	MANUAL_REVIEW
7	No contact route, DOM complete	FAIL — no mailto: , no link labelled or addressed as contact, support or helpdesk

Branches 2 and 3 are why depth 1 is worth the requests here. A "Contact" link is ambiguous on its own; the page behind it usually is not — it either names a helpdesk and offers a form or an address, or it does not. One request settles what keyword matching cannot.

Point 7 — the page is in English

[check_7](#). The only point using statistical detection rather than pattern matching.

Steps:

1. Take [main_text](#), falling back to [full_text](#). Read the declared [lang](#) attribute.
2. If under 120 characters → [MANUAL_REVIEW](#), too little to detect a language.
3. Build a [lingua](#) detector over all Latin-script languages, take the first 4000 characters, and record the top language and its confidence.

#	Detected	Declared	Verdict
1	en	en*	PASS

#	Detected	Declared	Verdict
2	en	anything else	PASS — the mismatch is a metadata inconsistency worth fixing, but it does not breach point 7
3	not en	en*	MANUAL_REVIEW — mixed-language pages and navigation-heavy text both cause this
4	not en	not en*	FAIL

If [lingua](#) is not installed, the check degrades rather than crashing: a declared [en](#) gives [PASS](#) noting that detection was unavailable, anything else gives [MANUAL_REVIEW](#).

Separating [main_text](#) from [full_text](#) matters most here. A cookie banner or a language switcher listing eight languages in the page chrome is exactly the kind of text that misleads a detector, and it is excluded by construction.

6. Where this workflow is known to be wrong

Listed because a checker whose limitations are undocumented invites more trust than it has earned.

Point 3 over-matched on the page's own hostname — fixed, and the first diagnosis of it was wrong. The haystack included [img.src](#) in full, so for a node served from a host containing "eosc" the regex matched the *hostname of every locally-hosted image*. The run of 21 September 2026 reported 8 "EOSC-referencing image assets" for that node; probing the captured evidence showed **7 of the 8 matched on the hostname alone** — a favicon counted three times, a parallax background, a gateway logo, a partner logo, and an EU funding badge named [FundedbytheEU.png](#). None is an EOSC logo. Exactly one was real.

This document previously recorded the cause as cross-host loading and proposed restricting the host. That was the wrong fix, and probing the evidence before applying it is what showed why: several nodes legitimately serve their EOSC lockup from a CDN or their own domain ([i0.wp.com](#), [research.csc.fi](#), [portal.eudat.eu](#), [www.panosc.eu](#)), so a host restriction would have discarded genuine logos. The defect was never cross-host loading; it was the node's own host being part of the text being searched.

The fix strips the host from [src](#) and requires "eosc" to be a token rather than a substring, with an explicit allowance for [eosc.eu](#)-hosted assets. Re-running against the same evidence moved that node from 8 assets to 1 and left all eight other nodes untouched. **No verdict changed** — the branch returns [MANUAL_REVIEW](#) either way — and the tally is still 28 PASS / 7 FAIL / 55 review; what changed is that the evidence line no longer pushes a reviewer towards the wrong conclusion.

The boundary rule then had a defect of its own, found by regenerating the report rather than by reasoning about it: requiring a non-word character *after* "eosc" rejected [EOSCNode_Finland-1-1-scaled.jpg](#), which is the official lockup, and cost the Finnish node its only EOSC asset — flipping its point 3 evidence from "asset present" to "none found". An uppercase letter is now accepted as a word break. The leading guard is still strict, which is what keeps [geoscience](#) and [GEOSCIENCE](#) out.

The approved-name match over-matched, then under-matched. Both are recorded because both shipped. First, the name was matched as a bare substring of the page body, so a short name matched inside ordinary words: [EGI](#) in the list reported a match on the BBMRI-ERIC page (inside "strategic") and on Data Terra (inside "Norwegian"), neither of which names the EGI node. The list was also matched as a whole, so any name matching any page satisfied that page's evidence line. Matching was made boundary-aware and scopable with a [node-id:](#) prefix.

That fix was then too strict. Run against the official Tripartite list, which writes every name as [EOSC Node | X](#), it matched **none of the nine pages**: the pages disagree about the separator glyph, not the name — BBMRI-ERIC writes a hyphen and an en dash, European DTO the pipe, EUDAT nothing at all. Publishing that would have been a finding against every node for a typographic difference. Separator glyphs are now interchangeable, and the evidence quotes what the page wrote so the variant is visible. This is a deliberate loosening, so it is reversible rather than baked in: [--strict-separators](#) restores the literal rule and reproduces the 0-of-9 result, and the report states which rule was in force either way. The place to challenge the default is the table in the guide.

The match is still not proof of the right name. The committed list is unscoped, so a match says the string is on the page, not that it is that page's own name; and with 9 names applying to all 9 nodes, a page carrying a *different* node's approved name would satisfy its own evidence line. The line says which claim it is making. Scoping the list is what makes the stronger claim available, and [checklist/approved-names-scoped.txt](#) now does. Run against the same evidence it returns the same two nodes, BBMRI-ERIC and EUDAT, which retires the concern for this dataset: the 2-of-9 headline does not depend on names leaking between pages. The published run of 24 September confirms it on twelve nodes: both lists match the same five (BBMRI-ERIC, Czechia, EUDAT, GÉANT, Slovakia), and [--strict-separators](#) matches none. The published run of 26 September gives the same agreement on thirteen: both lists match four (GÉANT's page was blocked that day), and [--strict-separators](#) matches none. The residual weakness is no longer the matching but the mapping — name to node id — which was derived in this repository rather than taken from the Tripartite file.

No verdict changed at any point, because the name never decides point 3 — but the evidence line did, twice, and both times it was wrong.

Nothing about visibility is captured, as set out in section 3. Point 3 cannot be closed automatically under the current evidence model, regardless of how the matching is fixed.

The vocabularies are English-only. A node serving a Dutch or Finnish contact page will be missed by points 5b, 5c and 6. Point 7 requires the landing page to be in English, so the bias is tolerable for this checklist — but it is a bias, not a neutral default.

A legitimate cross-host logo is indistinguishable from an unrelated one. At least one node in this set loads its node logo from [eossc.eu](#) while being served from its own domain, which is correct behaviour. A host check alone would need care not to reject it.

Binary policy documents are never read. [select_children](#) skips URLs ending in a binary extension, so an AUP published as a PDF is recorded as a pointer (branch 1 of [_policy_check](#)) and never verified. The checklist's requirement is accessibility, and the tool cannot confirm it for these.

--only used to narrow the run in place — fixed, with a cost. Assessing a subset rewrote the shared report, so a one-node table replaced the full one and looked complete. [--only](#) now narrows only which sites are fetched; the report written to the default results directory still covers every node, and the requested subset is recorded in a top-level [selection](#) list. An explicit [--results DIR](#) keeps the old narrowing, where nothing shared is at risk.

The cost is that such a report is no longer uniformly fresh: unselected rows come from evidence already on disk, under a header stating a single run timestamp. That is a freshness claim the report cannot make, so it no longer makes it — a **Mixed freshness** banner in both the Markdown and HTML output names what was re-fetched, counts what was not, and gives the capture date of the reused evidence. Nodes with no evidence at all are still listed under [missing_evidence](#), printed in red, with exit code 2.

The row heading comes from [nodes.yaml](#), not from the evidence. [assess](#) labels each node with the [url](#) configured now, while every verdict comes from the evidence captured then. After BBMRI-ERIC's URL changed on 25 September, a bare [assess](#) headed its row with the new [dev3](#). address above verdicts taken from the old [www](#). page. The address really fetched is kept only as [final_url](#), in the evidence and in the [fetch](#) summary of [results.json](#). Until the tool compares the two, rebuild an old run with the node list it was collected with (`--nodes`); the run guide gives the exact command.

A public CI summary is not the place for per-node verdicts. The workflow used to copy the whole of [results.md](#) into `$GITHUB_STEP_SUMMARY`, which on a public repository is world-readable. The verdicts are already public — the report is committed — so nothing leaked that was not already published; but that page is produced automatically from an unreviewed run and carries the repository's name, which makes a table of FAILs against named organisations read as a finding rather than a draft. The summary now publishes counts only, labelled as unreviewed, and points at the artifact for the per-node detail. It also refuses to summarise at all unless the check step succeeded, because [results.json](#) is committed and would otherwise have been read out of a fresh checkout as though it were this run's result.

7. Reading the output

[results.json](#) carries [run_id](#), [generated_at](#), the full [checklist](#) as loaded, [approved_names](#) (whether a list was used, whether it was the committed default ([default_used](#)), whether it was scoped, how many names, its path, the [sha256](#) of the bytes actually read, and [strict_separators](#) — plus the older [approved_names_supplied](#) flag, kept for readers of earlier result files), [selection](#) (the node ids `--only` restricted the fetch to, empty on a full run), [skipped](#) (the node ids `--skip` left out, empty when nothing was skipped), [url_mismatch](#) (present only when some node's evidence was collected from a URL other than the one now configured: [id](#), [configured_url](#), [evidence_url](#), [fetched_at](#); the reports then carry an "Evidence from a different URL" banner), and [nodes](#). Each node carries [id](#), [name](#), [url](#), [ad_hoc](#), a [fetch](#) summary (status, final URL, robots note, screenshot, crawl depth, every child with its [selected_for](#) and outcome, and [children_skipped](#)), [results](#), and [results_depth_1](#) when depth 2 was used.

The same run is rendered to [index.html](#) (the matrix, with evidence expandable per cell), [results.md](#), [results.csv](#), and [checklist-v3.0.html](#) (named after the revision applied, so [checklist-v3.1.html](#) after a switch; the checklist itself, so a reader can see the rule a verdict was derived from without opening the source document).

The honest summary of the whole workflow: of the ten points, **two** (1 and 4) are routinely settled by inspection, **three** (5b, 5c, 6) are settled when depth 1 gives the tool a second page to read, and **five** (1R, 2, 3, 5a, and 7 in its mixed-language branch) end in a human's hands by design. A tool that returned PASS or FAIL on every line would look more useful and be worth considerably less.

8. Worked examples

Three step-by-step procedures, one for each change that comes up most often: adding a node, changing a node's URL, and adopting a new revision of the checklist. Each lists which files to edit, the exact commands, and how to put the result into the published [results/](#) directory.

Four rules apply to all three:

- **Try it in a scratch folder first.** `--results /tmp/trial` writes nowhere near `results/`, so nothing published changes until you decide it should.
- **Run `uv run pytest -q` after every edit.** It needs no network, takes a few seconds, and names the file you got wrong.
- **Contact as few sites as possible.** Only `collect` and `run` send requests. `assess` works on saved evidence, so it can be repeated as often as you like.
- **A new published run needs a human review before it is committed.** The tool's verdicts are not a compliance statement (see "Reading the output" in the [run guide](#)). Write the review down, as [results/REVIEW-2026-09-24.md](#) does for the 24 September run, and commit only then.

In terms of the three stages in section 1: examples 1 and 2 need **stage 1** (collection) again, but only for the one node concerned. Example 3 changes **stage 2** (the rules), so it can be tried on the evidence already collected, without contacting any site. Stage 3 (the reports) is always regenerated by `assess`.

The run labels below (`live-2026-10-01-...`) are examples. Use the date of your own run.

Example 1 — adding a node and publishing it

The example adds a node with the placeholder id `eosc-example`. Replace every value with the real ones.

Step 1: edit three files. All three are at the repository root or in `checklist/`, and all are edited by hand.

`nodes.yaml`: add the entry. `url` is the landing page as registered in the EOSC EU Node Contributors Dashboard ("Website address", section 1.2, field 6). Copy `eosc_page` from the [federation index](#); do not guess it.

```
- id: eosc-example
  name: EOSC Node Example
  url: https://eosc-node.example.org/
  eosc_page: https://eosc.eu/building-the-eosc-federation/eosc-node-example/
```

`checklist/approved-names.txt`: add the approved name exactly as the Tripartite list writes it:

```
| EOSC Node | Example
```

`checklist/approved-names-scoped.txt`: add the same name, with the node id in front:

```
| eosc-example: EOSC Node | Example
```

If the Tripartite list does not include the node yet, do not make a name up. Leave both name files alone. `test_nodes.py::test_every_node_has_a_scoped_approved_name` then fails and names the node, which is intended: the node is not ready to be published. In the meantime, check it in a scratch folder with `--url https://eosc-node.example.org/` (see "Checking a page not in `nodes.yaml`" in the [run guide](#)), which needs no configuration.

Step 2: check the configuration, offline.

```
| uv run pytest -q
```

A missing scoped name, an id that is not a valid filename, a URL that is not `https`, an `eosc_page` outside `eosc.eu`, or a landing page already used by another node each fail with a message naming the problem.

Step 3: trial run for the new node only. This contacts one site.

```
uv run basic-check collect --only eossc-example --results /tmp/trial
uv run basic-check assess --only eossc-example --results /tmp/trial
uv run basic-check show eossc-example --results /tmp/trial
```

Open `/tmp/trial/index.html` or `/tmp/trial/results.md` and read the new row. If the page returned `ERROR` (robots.txt, a 403, a DNS failure), stop there. An `ERROR` row is not a result worth publishing.

Step 4: put the node into the published run. There are two ways.

A — *add the one node to the current run (one site contacted, and already done).* Reuse the evidence from step 3 rather than fetching the page again:

```
cp /tmp/trial/evidence/eossc-example.json results/evidence/
cp /tmp/trial/evidence/screenshots/eossc-example.png results/evidence/screenshots/
uv run basic-check assess --only eossc-example --skip Italy \
--run live-2026-10-01-plus-eossc-example
```

The report covers every node. A **Mixed freshness** banner says that only `eossc-example` is new and that the other rows are reused from the earlier capture, with its date. `--skip Italy` keeps the same scope as the published run. Italy's URL changed on 26 September to <https://eoscnod-it.d4science.org/>; drop `--skip Italy` once a trial collection (`collect --only eossc-it --results /tmp/trial`) shows the page loads.

B — *a complete new run (every node contacted once).* Use this when the existing evidence is old enough that a single timestamp is worth more than leaving the other sites alone:

```
uv run basic-check run --skip Italy --results /tmp/new-run --run live-2026-10-01
# review /tmp/new-run/results.md and index.html, then:
rm -rf results/evidence
cp -r /tmp/new-run/evidence results/evidence
uv run basic-check assess --skip Italy --run live-2026-10-01
```

The last command re-renders the reports in `results/` from the evidence you reviewed. It makes no requests.

Either way, `assess` must end with **exit code 0**. Exit 2 means a configured node has no evidence, so the table is incomplete.

If `assess` warns "**Evidence from a different URL**". It means a node's URL in `nodes.yaml` is not the page its evidence was collected from. That is the case for BBMRI-ERIC from 25 to 26 September 2026, when `nodes.yaml` had the `dev3`. address but the published evidence was from the `www`. one. The warning also appears as a banner in both reports. Either collect that node again first (example 2), or assess with the node list the evidence was collected with, plus the new entry:

```
git show 53081f6:nodes.yaml > /tmp/nodes-published.yaml
# append the eossc-example entry from step 1 to /tmp/nodes-published.yaml
uv run basic-check assess --nodes /tmp/nodes-published.yaml --only eossc-example \
--skip Italy --run live-2026-10-01-plus-eossc-example
```

Step 5: re-run the tests against the new published run.

```
uv run pytest -q
```

Two tests read the committed evidence. Each exists to make a change visible, not to be satisfied mechanically:

- `test_privacy.py::test_the_committed_evidence_publishes_no_personal_address_or_phone` must pass as it stands. If it fails, personal data is about to be published. Stop and find out why; never change the test to let it through.
- `test_names.py::test_the_official_names_match_the_nodes_that_show_them` pins which nodes display their approved name (four for the 26 September run; five on 24 September, when GÉANT's page was served). If the new page shows its name, add its id to the expected set and update the docstring to say why.
- `test_update_row.py::test_published_results_are_the_run_these_tests_copy` pins how many nodes the run has and which were skipped (13 and none, for run `web-4`). Update it to the new run's shape.

Step 6: update what quotes the published figures, then commit. `results/` is regenerated, but these are written by hand:

- a review document for the new run, like `results/REVIEW-2026-09-24.md`;
- the headline tally in `README.md` ("130 cells: 45 PASS · 8 FAIL · 77 review");
- "The published figures" in `docs/TEST-SUITE.md`, and the node list in "The node list — a YAML file" in the same file.

```
git status
git add nodes.yaml checklist/approved-names.txt checklist/approved-names-scoped.txt \
    results README.md docs tests
git commit -m "Add eoscc-example; publish run live-2026-10-01-plus-eoscc-example"
git push
```

Example 2 — changing a node's landing page URL

The example is the change actually made on 25 September 2026: BBMRI-ERIC moved from <https://www.bbmri-eric.eu/eoscc-node-bbmri-eric/> to <https://dev3.bbmri-eric.eu/eoscc-node-bbmri-eric/>. It was set back on 26 September, because the `dev3.` host's `robots.txt` excludes the checker, so `nodes.yaml` has the `www.` address again.

Which files change. One line in one file: `url` for that node in `nodes.yaml`. Nothing else in the repository stores a node's URL. Add a comment recording the old address and the date, so the published evidence can still be traced. `--update-nlp` makes both edits and changes nothing else (it contacts no site and commits nothing):

```
uv run basic-check --update-nlp bbmri-eric https://new.example.org/eoscc-node/
```

To go back, `uv run basic-check --restore-default-config` puts `nodes.yaml` back as it was downloaded from GitHub and keeps the replaced file as `nodes.yaml.<date-time>.bak`.

Done by hand, the entry looks like this:

```
- id: bbmri-eric
  name: BBMRI-ERIC
  # URL changed on 25 September 2026, from
  # https://www.bbmri-eric.eu/eoscc-node-bbmri-eric/
  url: https://dev3.bbmri-eric.eu/eoscc-node-bbmri-eric/
  eoscc_page: https://eoscc.eu/building-the-eoscc-federation/eoscc-node-bbmri-eric/
```

Change `eoscc_page` as well only if the node's own entry on `eoscc.eu` moved; a new landing page does not usually mean a new `eoscc.eu` entry. The name files change only if the node's approved name changed. Keep the `id`: it names the evidence file, and changing it would make the node look new.

Step 1: check the configuration.

```
uv run pytest -q
```

Step 2: trial run for that node only. This contacts one site.

```
uv run basic-check collect --only bbmri-eric --results /tmp/trial
uv run basic-check assess --only bbmri-eric --results /tmp/trial
uv run basic-check show bbmri-eric --results /tmp/trial
```

Check that the page actually loaded before going further. On 25 September the `dev3.` host's `robots.txt` disallowed every path, the tool honoured it, and all ten points were `ERROR`. A result like that is not worth publishing: ask the node for its public address, or for the checker to be allowed, and stop there.

Step 3: update that node's row in the published run. Only BBMRI-ERIC is fetched again; the other rows are copied from `results/results.json` as they are, neither re-fetched nor re-judged:

```
uv run basic-check run --update-results-for-node bbmri-eric
```

That fetches the one page into a scratch directory, and only if the page loaded does it replace `bbmri-eric`'s evidence, screenshot and row, then rebuild `results.md`, `index.html` and `results.csv` from the stored rows. A failed fetch changes nothing (add `--accept-error` to record it anyway). It prints each point's old and new verdict, and both reports gain an **Updated rows** banner naming the node, the date and the previous URL. The run id and time in the header stay those of the original run. To reuse the evidence from step 2 instead of fetching again, copy it into `results/evidence/` and run `uv run basic-check assess --update-results-for-node bbmri-eric`, which is offline. To see the outcome before touching the published run, try it on a copy first: `cp -r results /tmp/copy` and add `--results /tmp/copy`. For a completely fresh table instead, follow path B of example 1 step 4.

Why the order matters. `assess` labels each row with the URL in `nodes.yaml`, but takes the verdicts from the evidence on disk. Between step 1 and step 3, the two describe different pages. `assess` now detects this. It prints a warning naming the node and both URLs, puts an **Evidence from a different URL** banner in `results.md` and `index.html`, and records the pair under `url_mismatch` in `results.json`. Step 3 clears it, because the evidence now comes from the configured URL.

To rebuild a run collected with an older node list, without the warning, use that list. For the reviewed run of 24 September (in the git history at `800d632`), assessed in a scratch copy of its `results/`:

```
git show 53081f6:nodes.yaml > /tmp/nodes-2026-09-24.yaml
uv run basic-check assess --run live-2026-09-24-no-italy --skip Italy \
--nodes /tmp/nodes-2026-09-24.yaml
```

Step 4: tests, review, commit. As in example 1, steps 5 and 6: run `uv run pytest -q`, check the two tests that read the committed evidence, write the review, update the quoted figures, and commit `nodes.yaml`, `results/` and the documents together.

Example 3 — moving to a new revision of the checklist

There is no URL to change. The tool never downloads the checklist document. A copy of it is committed in `checklist/`, and three things connect the tool to it:

Where	What it says
<code>src/basic_check/cli.py</code> , the line <code>DEFAULT_CHECKLIST = ROOT / "checklist" / "v3.0.yaml"</code>	Which checklist revision is used when <code>--checklist</code> is not given. This is the one line to change to make a new revision the default. The help texts, the tests and the reports all follow it.

Where	What it says
checklist/vX.Y.yaml: checklist_version, checklist_date, source_document, source_file, source_sha256	Which document that revision was transcribed from, and its SHA-256 hash.
--checklist / -c <file> on assess, run and points	A different revision for a single run, without changing the default.

When the originating document changes, you therefore add a new file for the new revision and point the default at it. `v3.0.yaml` is never edited: the published reports name the file they were produced from.

The example assumes a v3.1 dated 15 October 2026. Use the real version, date and filenames.

Step 1: commit the new document next to the old one.

```
cp ~/Downloads/20261015_Node_Landing_Page_Verification_Checklist_v3.1.pdf checklist/
sha256sum checklist/20261015_Node_Landing_Page_Verification_Checklist_v3.1.pdf
# on macOS: shasum -a 256 checklist/20261015_Node_Landing_Page_Verification_Checklist_v3.1.pdf
```

The committed v3.0 file is a PDF rendering of the circulated `.docx`. Commit the `.docx` itself if you prefer: `source_file` may name either.

Step 2: create the new revision's file from the old one.

```
cp checklist/v3.0.yaml checklist/v3.1.yaml
```

In `checklist/v3.1.yaml`, change the header:

```
checklist_version: "3.1"
checklist_date: "2026-10-15"
source_document: 20261015_Node_Landing_Page_Verification_Checklist_v3.1.docx
source_file: 20261015_Node_Landing_Page_Verification_Checklist_v3.1.pdf
source_sha256: <the hash printed in step 1>
```

The version must match the filename (`v3.1.yaml` declares "3.1"); a test enforces this.

Step 3: compare every point with the new document. No tool can do this step for you. For each point in `v3.1.yaml`:

- If the wording changed, update `title` and `requirement`, which are quoted word for word. Then read the `check_*` function named in `implemented_by`, in `src/basic_check/checks.py`, and change it if it no longer applies the new wording.
- If what a script can decide changed, update `decidable` and `decidable_note`.
- Where a rule changed, update that point's decision procedure in section 5, so the written procedure and the code still describe the same rule.
- If a point was **added**, write a new `check_<id>` function in `checks.py`, call it from `run_all()` in `checklist` order, name it in the point's `implemented_by`, and add tests for it to `tests/test_checks.py`.
- If a point was **removed**, delete its entry, its function and its tests.

Step 4: try the new revision offline, before switching. `-c` applies it to one run only. Re-scoring saved evidence contacts no one:

```
uv run basic-check points -c checklist/v3.1.yaml
mkdir -p /tmp/v31 && cp -r results/evidence /tmp/v31/
```

```
git show 53081f6:nodes.yaml > /tmp/nodes-2026-09-24.yaml
uv run basic-check assess -c checklist/v3.1.yaml --results /tmp/v31 \
--nodes /tmp/nodes-2026-09-24.yaml --skip Italy --run trial-v3.1
```

Compare `/tmp/v31/results.md` with `results/results.md`. Any difference comes from the new rules alone, since the evidence is the same.

Step 5: make it the default, and run the tests. In `src/basic_check/cli.py`:

```
DEFAULT_CHECKLIST = ROOT / "checklist" / "v3.1.yaml"
uv run pytest -q
uv run basic-check --help # now says "against checklist v3.1"
```

The tests check the new file's hash, that its version matches its filename, and that every point maps to exactly one `check_*` function and every function to one point. They also check that `v3.0.yaml` and its document are still intact, so older runs can still be rebuilt with `-c checklist/v3.0.yaml`. What they cannot check is whether a rule is right: that is step 3.

Step 6: produce and publish the results. The default revision is now v3.1, so no `-c` is needed. Choose between the same two paths as in example 1:

```
# fresh evidence from every node:
uv run basic-check run --skip Italy --results /tmp/new-run --run live-2026-10-20-v3.1
# review, then:
rm -rf results/evidence && cp -r /tmp/new-run/evidence results/evidence
uv run basic-check assess --skip Italy --run live-2026-10-20-v3.1
```

Alternatively, re-score the published evidence against v3.1 without contacting anyone, as in step 4 but writing to `results/`. The report then says v3.1, while the capture date is still that of the evidence.

The reports now link to `results/checklist-v3.1.html`. The old `results/checklist-v3.0.html` stays behind; remove it, and update the places that name v3.0 in prose or links:

```
git rm results/checklist-v3.0.html
rg -n "v3\0" README.md docs checklist/README.md .github
```

In `README.md`, the table of checklist points links each point to `results/checklist-v3.0.html#p....`. Those links break once the file is removed, so change them to v3.1.

Step 7: tests, review, commit. As in example 1, steps 5 and 6. If any `check_*` function changed in step 3, also update its decision procedure in section 5 of this document, which describes each point's rule branch by branch. Commit the new document, `checklist/v3.1.yaml`, `cli.py`, any changed checks and tests, `results/` and the documents together, so the commit shows the whole change.

Source checklist: Node Landing Page Verification Checklist v3.0, 15 September 2026, transcribed to `checklist/v3.0.yaml` alongside the source document. Implementation: `src/basic_check/{fetch,patterns,checks,names,privacy,report,cli}.py`. Repository: <https://github.com/marioreale/eosc-basic-compliance>.