

eosc-basic-compliance

Download, installation, configuration and run guide

Verified against a clean clone on 26 September 2026 — documentation commit [ee12687](#)

Repository: github.com/mario reale/eosc-basic-compliance

Contents

1. Requirements · 2. Installation · 3. Configuration · 4. Running it · 5. How deep to go: [--depth](#) · 6. Reading the output · 7. The test suite · 8. Continuous integration · 9. Troubleshooting · 10. Extending it · 11. Worked examples

[eosc-basic-compliance](#) checks EOSC Node Landing Pages against the **Node Landing Page Verification Checklist v3.1** (24 September 2026) and writes a report: a matrix of nodes against checklist points, with the evidence behind every verdict.

This guide covers installing it, configuring it for your own nodes, running it, reading what comes out, and running the test suite. This edition was checked against a clean clone of commit [ada1b4a](#) on 25 September 2026, and the test count again at commit [7ed46dd](#) on 26 September 2026. Every command that needs no network — installation, the test suite, [points](#), [assess](#) and the rebuild of the published report — was re-run that day, and the test counts, disk sizes and verdict tallies were re-measured rather than carried over. No node website was contacted to prepare it. Request counts and HTTP behaviour come from the runs of 24 and 26 September and from a trial run on 25 September, and each figure says which. Where a command's behaviour is surprising, that is noted rather than smoothed over.

The published run changed on 26 September 2026. [results/](#) now holds run [web-5](#), collected by the GitHub workflow on 26 September from all thirteen nodes under checklist v3.1 (45 PASS, 8 FAIL, 77 MANUAL_REVIEW), after run [web-4](#) of the same morning, whose verdicts it repeats cell for cell. It has **not** been reviewed by hand; the review of the run it replaced stays in [results/REVIEW-2026-09-24.md](#) as history. Passages below that describe "the published run of 24 September" are about that earlier, reviewed run, whose files are in the git history at commit [800d632](#).

Checklist v3.1 since 26 September 2026. The default reference checklist is now v3.1 of 24 September 2026 ([checklist/v3.1.yaml](#)). The evidence of [web-4](#) was first re-scored against it offline, then [web-5](#) was collected under it; every verdict is unchanged. The header of [checklist/v3.1.yaml](#) lists the differences point by point.

What this tool will not do: it does not produce a compliance statement. Of the ten checklist points, three can be settled by inspection, four only partly, and three require a human reading the page. The tool returns [MANUAL_REVIEW](#) rather than guessing, and roughly 60% of cells land there by design. A checker that returned PASS/FAIL on every line would look more useful and be worth less.

👉 **Analysis workflow** — the companion to this guide. This one tells you how to run the tool; that one tells you what it does once running: the collection sequence, the evidence model, and for each of the ten checklist points the exact branch conditions and the verdict each one yields. Read it when you want to predict a verdict, argue with one, or change a check.

1. Requirements

Python	3.12 or newer (the project is developed on 3.14)
Package manager	uv
Browser	Chromium, via Playwright — only needed to collect evidence
Disk	~500 MB for the virtualenv, ~660 MB more for Chromium (see the note below)
Network	Outbound HTTPS to the node websites, for collection only

Installing `uv`, if you do not have it:

```
| curl -LsSf https://astral.sh/uv/install.sh | sh # macOS / Linux
```

You do not need to create a virtualenv or run `pip install`. `uv` manages the environment from `uv.lock`, so everyone gets identical dependency versions.

On the disk figures. Both were measured with `du -sh`, not estimated. The 500 MB is `.venv` in the project directory (484 MB in a fresh clone on 25 September 2026, rising to about 510 MB once caches accumulate). The 660 MB is one Playwright install — Chromium (393 MB), the headless shell (261 MB) and ffmpeg (5 MB) — and it does **not** live in the project: Playwright puts browsers in a shared cache at `~/.cache/ms-playwright`, so a second project on the same machine reuses it and costs nothing. That cache can nonetheless grow well past 660 MB, because each Playwright version keeps its own browser build; the machine this was measured on held two versions and so 1.3 GB in total. Delete the whole directory and re-run `playwright install` to reclaim it.

2. Installation

```
| git clone https://github.com/marioreale/eosc-basic-compliance.git
| cd eosc-basic-compliance
| uv sync --locked
```

This creates `.venv/` and installs exactly the versions pinned in the committed `uv.lock`: it resolves 27 packages and installs 26, the 27th being the project itself. Expect it to take under a minute; from a warm `uv` cache it took two seconds. A plain `uv sync` also works, but it will silently re-resolve if the lockfile and `pyproject.toml` ever disagree; `--locked` fails instead, which is what you want when the point is to reproduce someone else's result. CI uses `--locked` for the same reason (section 8).

At this point you can already run the test suite, read the checklist, and regenerate the committed report from the evidence in the repository. To **fetch pages**, you also need a browser:

```
| uv run playwright install chromium --with-deps
```

This downloads roughly 660 MB (see the note in section 1). On Linux, `--with-deps` also installs system libraries and will ask for `sudo`. If you only want to re-render reports from evidence already in the repository, skip this step.

What needs a browser, and what does not

This distinction saves a large download in CI and on review machines:

Command	Browser	Network	What it does
points	no	no	Prints the checklist and what is machine-decidable
assess	no	no	Applies the checklist to evidence already on disk
show	no	no	Prints one node's results in the terminal
collect	yes	yes	Fetches pages and saves evidence
run	yes	yes	collect , then assess
pytest	no	no	The whole test suite is offline

Verified: with `PLAYWRIGHT_BROWSERS_PATH` pointed at an empty directory, all 426 tests still pass, while `collect` fails with Playwright's Executable doesn't exist ... [run playwright install](#).

Verifying the installation

```
uv run pytest -q           # expect: 426 passed
uv run ruff check src tests # expect: All checks passed!
uv run basic-check points  # prints the ten checklist points
```

If `pytest` passes, the checklist transcription, the check logic and the report renderer are all working. No network access is involved.

3. Configuration

Two files control what is checked, and a third optional one supplies an input the repository cannot ship. The two are plain YAML and both are meant to be edited; the third is a plain text list you write yourself.

`nodes.yaml` — which pages to check

```
nodes:
- id: bbmri-eric
  name: BBMRI-ERIC
  url: https://www.bbmri-eric.eu/eosc-node-bbmri-eric/
  eosc_page: https://eosc.eu/building-the-eosc-federation/eosc-node-bbmri-eric/
```

Field	Meaning
id	Short slug. Used for evidence filenames and for <code>--only</code> / <code>show</code> .
name	Display name in the report.
url	The URL registered in the EOSC EU Node Contributors Dashboard ("Website address", section 1.2 field 6). That registered URL is what the checklist defines as the Node Landing Page. If the URL here differs from the registered value, the registered value wins — correct this file.

Field	Meaning
<code>eosc_page</code>	The node's dedicated entry on eosc.eu , which checklist point 4 requires the landing page to link to.

`eosc_page` is what lets a point 4 failure name the exact URL that is missing, rather than reporting a bare absence. The slugs in the committed file were read from the live index at eosc.eu/building-the-eosc-federation/. If you add a node, look its slug up there rather than guessing it from the name.

BBMRI-ERIC's URL was changed and then set back. On 25 September 2026 it was changed to <https://dev3.bbmri-eric.eu/eosc-node-bbmri-eric/>; on 26 September it was set back to the [www](https://www.bbmri-eric.eu/eosc-node-bbmri-eric/). address shown above, because the `dev3`. host's `robots.txt` excludes every path. `nodes.yaml` carries a comment saying so. The published run in [results/](#) (26 September) was collected from this [www](https://www.bbmri-eric.eu/eosc-node-bbmri-eric/). address, as was the reviewed run of 24 September, so the configuration and the published evidence agree.

EOSC Node Italy's URL changed on 26 September 2026, from <https://eosc.it/> to <https://eoscnode-it.d4science.org/>, and `nodes.yaml` carries a comment saying so. eosc.it had no address record on 24 September, so that run skipped Italy; nothing published depends on the old address. Italy was assessed for the first time in run [web-4](#) on 26 September, from the new address, and again in the published run [web-5](#).

Changing a node's URL

Change `url` in `nodes.yaml` and run `uv run pytest -q`. Nothing else in the repository stores the URL. You can edit the file by hand, or let the tool do it:

```
uv run basic-check --update-nlp bbmri-eric https://new.example.org/eosc-node/
```

`--update-nlp` takes the node id (the id exactly, not the name) and the new URL. It changes only that node's `url`: line, adds a comment above it with the date and the old URL, and keeps every other line and comment as it was. It refuses a URL that is not absolute https, one that is already another node's landing page, and an unknown id, and it prints the old and new URL. It contacts no site, commits nothing and leaves [results/](#) alone; check the change with `git diff nodes.yaml` and commit it yourself.

To undo it, and any other edit to `nodes.yaml`, put the file back as it was downloaded from GitHub:

```
uv run basic-check --restore-default-config
```

The replaced file is kept beside it as `nodes.yaml.<date-time>.bak` (ignored by git), and the output lists every node whose URL was restored, added back or removed. If the file is already the default, nothing changes. The default is a copy shipped inside the package (`src/basic_check/defaults/nodes.yaml`), so this works without git, from a zip download too. Only `nodes.yaml` is restored: the approved-name lists in [checklist/](#) are not touched. Like `--update-nlp`, it contacts no site, commits nothing and leaves [results/](#) alone; for a node whose URL it restored, run `--update-results-for-node <id>` re-checks that node. To re-check only that node and update only its row in the results, keeping the others, use `run --update-results-for-node` (section 4). To look first without changing anything published, collect it into a scratch directory first, so you can see what the new page gives before anything published changes:

```
uv run basic-check collect --only bbmri-eric --results /tmp/trial
uv run basic-check assess --only bbmri-eric --results /tmp/trial
```

`assess` takes each node's URL from `nodes.yaml`, not from the evidence. So after a URL change, a bare `uv run basic-check assess` pairs the new URL with evidence collected from the old page. It no longer does so silently. `assess` compares each node's configured URL with the `requested_url` stored in its evidence. When

they differ, it prints a warning naming the node and both addresses, adds an **Evidence from a different URL** banner to [results.md](#) and [index.html](#), and lists the pairs under `url_mismatch` in [results.json](#). The exit code does not change, because the table is complete; it is just not about the configured page. The address finally reached, after redirects, is kept as `final_url` in the evidence file.

To rebuild the published run, assess it under its own label. It was collected from the committed [nodes.yaml](#), so no other option is needed and there is no URL warning:

```
uv run basic-check assess --run web-5
```

Verified on 26 September 2026: this reproduces the committed [results/](#) exactly, apart from the generation time, with 45 PASS, 8 FAIL and 77 MANUAL_REVIEW. The reviewed run of 24 September that it replaced is in the git history at commit [800d632](#); to rebuild that one, restore its [results/](#) into a scratch directory and assess it with `--skip Italy` and the node list from commit [53081f6](#). The full procedure, including how to publish the new node's results, is [example 2](#) in section 11.

A new URL can behave differently from the old one in ways that have nothing to do with the checklist. On the 25 September trial, [dev3.bbmri-eric.eu](#) served a [robots.txt](#) of `User-agent: * / Disallow: /`. The tool honours it, so no page was requested and all ten points for BBMRI-ERIC are **ERROR** (section 6). That is typical of a staging host and says nothing about compliance. Ask the node for the public address, or for the development host to allow the checker, before the next published run.

checklist/v3.1.yaml — the rules

One entry per checklist point. The `requirement` text is quoted verbatim from the source document so a reader can audit the rule without a diff, and `decidable` records whether a script can settle the point at all:

- `true` — decidable by inspection (points 1, 4, 7)
- `partial` — partly decidable; the tool gathers evidence and defers (3, 5b, 5c, 6)
- `false` — needs a human reading the page (1R, 2, 5a)

The file also pins its own provenance: `source_file` names the committed PDF the transcription was made from and `source_sha256` is that file's hash, asserted by [tests/test_checklist.py](#). Be clear about what this proves — it pins the exact bytes the transcription came from, which catches silent drift when a checklist is revised. It is not a signature, and the committed PDF is a rendering of the circulated `.docx`, not that `.docx` itself.

The checklist is not downloaded from anywhere. No file in the repository holds a URL for it. The document is committed in [checklist/](#), and one line in [src/basic_check/cli.py](#) selects the revision used by default:

```
DEFAULT_CHECKLIST = ROOT / "checklist" / "v3.1.yaml"
```

`--checklist / -c` overrides it for a single run. The help texts, the report file name ([checklist-v3.1.html](#)) and the tests all follow that line.

If the checklist is revised to v3.2, add a new [checklist/v3.2.yaml](#) with its own source file and hash rather than editing v3.1 in place, then change that line. Past runs should remain reproducible against the rules that produced them. The full procedure is [example 3](#) in section 11.

The approved-names file — `--approved-names`

Checklist point 3 has two halves: an EOSC logo, and **the official Tripartite-approved node name**. The official list of those names is committed to the repository at [checklist/approved-names.txt](#), and is used automatically. You do not have to supply anything for the name half to be assessed.

EOSC Node	BBMRI-ERIC
EOSC Node	European DTO
EOSC Node	Data Terra
EOSC Node	Finland
EOSC Node	PaNOSC
EOSC Node	EUDAT
EOSC Node	EGI
EOSC Node	GÉANT
EOSC Node	EBRAINS RI

Which list a run used is recorded in the report, in one line at the top of both the HTML and the Markdown, and in `results.json` under `approved_names`. A reader should never have to assume that the run they are looking at checked the names they think it did.

Which list is used

You run	List used
<code>basic-check assess</code>	the committed <code>checklist/approved-names.txt</code>
<code>basic-check assess --approved-names mine.txt</code>	<code>mine.txt</code> , instead of the default
<code>basic-check assess --no-approved-names</code>	none; point 3's name half is not assessed

A `--approved-names` path that does not exist is a clean error, **not** a silent fall back to the default. You named a specific list; checking a different one is not a smaller failure than checking none.

Both flags work identically on `assess` and on `run`.

Supplying your own list

One entry per line. Blank lines and `#` comments are ignored; write `\#` for a literal hash in a name.

```
# My own list, 21 September 2026.
bbmri-eric: EOSC Node - BBMRI-ERIC # applies to that node only
EGI Node                          # applies to every node
```

Prefix each name with the node id from `nodes.yaml`. That ties the name to the node it belongs to, and it is the difference between the tool reporting "this page shows this node's approved name" and the much weaker "some approved name appears somewhere on this page". The committed official list is *not* scoped — it arrived as bare names, with nothing tying a name to a node — so runs using the default make the weaker claim, and say so. If you need the stronger one, copy the file, add the `node-id:` prefixes, and pass it with `--approved-names`.

```
uv run basic-check assess           # official list
uv run basic-check assess --approved-names my-names.txt # your list
uv run basic-check assess --no-approved-names           # no list at all
```

A file called `approved-names*.txt` in the repository root is gitignored, so a list you drop next to the project to override the default does not reach the repository by accident. The committed default, being under `checklist/`, is unaffected.

How a name is matched

Behaviour	Consequence
Matching is case-insensitive, and each word is literal	Regular-expression characters in a name are matched as themselves.

Behaviour	Consequence
A name must not be part of a longer word	EGI matches "the EGI Foundation" but not "strategic". A hyphen counts as part of a word, so BBMRI does not match BBMRI-ERIC .
The separator between words is flexible	Whitespace and the glyphs pipe, hyphen, en dash, em dash, colon, slash and middle dot are interchangeable — see below. A name that wraps across two lines still matches.
Punctuation inside a word is still required	BBMRI-ERIC does not match a page writing BBMRI ERIC .
The first match in file order is reported	Put the preferred form first if a node has more than one approved name.
A variant rendering is reported	The evidence quotes what the page actually writes, so you can see that it differs from the official form without diffing by eye.
The name never changes the verdict	Point 3 stays MANUAL_REVIEW either way. The list adds an evidence line for the reviewer; it cannot produce a PASS .
A page that was never fetched is never reported as missing the name	In the 21 September run GÉANT answered HTTP 403 with no body; the evidence said the name was not looked for, not that it is absent.

All of these match a list entry written [EOSC Node | BBMRI-ERIC](#):

[EOSC Node | BBMRI-ERIC](#)
[EOSC Node - BBMRI-ERIC](#)
[EOSC Node – BBMRI-ERIC](#)
[EOSC Node: BBMRI-ERIC](#)
[EOSC Node BBMRI-ERIC](#)

Only the body text is searched. The match runs against [full_text](#), which excludes the HTML `<title>`. The EGI page's title is [EGI Node - EGI](#), and that string appears nowhere in the body, so a name taken from the browser tab will not match. Take the name from the visible page.

Two of the matching rules above are fixes for defects, not features, and are described here so you can judge the results rather than trust them:

- **Word boundaries.** An earlier version matched a bare substring, so [EGI](#) in the list reported a match on the [BBMRI-ERIC](#) page (inside "strategic") and on Data Terra (inside "Norwegian"). Neither page names the EGI node.
- **Flexible separators.** The official list writes every name as [EOSC Node | X](#). The pages do not: [BBMRI-ERIC](#) writes a hyphen and an en dash, [European DTO](#) writes the pipe, [EUDAT](#) writes no separator at all. Matching the pipe literally scored **0 of 9** against the real pages — a finding against every node for what is a typographic difference. Treating the glyph as interchangeable is a deliberate loosening, and it is the reason the evidence line quotes what the page actually writes. Because it is a judgement call and not a fact, it is reversible from the command line: `--strict-separators` restores literal glyph matching, and the report then states that the strict rule was in force rather than the flexible one. Against the evidence of 21 September 2026 the strict rule matches **none** of the nine nodes, which is

the 0-of-9 result above. Against the published evidence of 24 September it matches none of the twelve.

A match is not proof of the right name

The default list is **unscoped**: any name in it, matching anywhere in a page's body, counts for any node. So a match tells you an approved name appears on the page — not that it is *that* page's own name. A node listing its partners would match on a partner's name.

`checklist/approved-names-scoped.txt` closes that gap. It writes each name as `node-id: EOSC Node | X`, so a name is only ever matched against the node it belongs to:

```
uv run basic-check assess --approved-names checklist/approved-names-scoped.txt
```

Running it against the published evidence of 24 September 2026 returns **the same five nodes** as the unscoped list: BBMRI-ERIC, Czechia, EUDAT, GÉANT and Slovakia. On the evidence of 21 September it returned the same two, BBMRI-ERIC and EUDAT. That is a useful result rather than a null one: the headline is not an artefact of names leaking across pages, because tying each name to its own node changes nothing.

It is **not** the default, and the reason is worth stating plainly: the *names* in it are the official ones, but the *mapping* from each name to a node id in `nodes.yaml` was derived in this repository and is not part of the Tripartite file. That mapping is the line to challenge before relying on a scoped run. The file's header comment says the same thing.

Which list was used is recorded

A report that cites a name list is only as trustworthy as your ability to tell which bytes it read. Every run records the SHA-256 of the list in `results.json` under `approved_names.sha256`, and the report prints the first twelve characters of it beside the file name. If the official list is revised and a run was made against the old one, the digests differ and you can see it without re-running anything.

```
sha256sum checklist/approved-names.txt
```

What the official list actually finds

In the published run of 26 September 2026, the committed list of thirteen names matches **four of the thirteen** nodes assessed. Every one of the four writes the name with a different separator from the list, so none of them would match literally:

Node	Outcome
BBMRI-ERIC	matched; the page writes <code>EOSC Node - BBMRI-ERIC</code>
EOSC Node Czechia	matched; the page writes <code>EOSC Node Czechia</code>
EUDAT	matched; the page writes <code>EOSC Node EUDAT</code>
EOSC Node Slovakia	matched; the page writes <code>EOSC Node Slovakia</code>
European DTO, Data Terra, PaNOSC	no match; the phrase <code>EOSC Node</code> occurs (5, 2 and 2 times), but never followed by an approved name
CERN, EOSC Finland, EGI, EBRAINS, Italy	no match, and the phrase <code>EOSC Node</code> does not occur in the body text
GÉANT	not looked for: the page answered HTTP 403 with the Cloudflare challenge

CERN's configured page is a sign-in endpoint with very little text, so there is almost nothing to match against. Italy's page shows "EOSC Node | Italy" in its logo image, which is not text. GÉANT's page, when served on 24 September, wrote [EOSC Node GÉANT](#) and matched, which made five of twelve that day.

The earlier run of 21 September matched two of nine (BBMRI-ERIC and EUDAT). GÉANT was not looked for then, because it answered HTTP 403 with no body, and CERN, Czechia and Slovakia were not yet configured.

This is a finding to review, not a verdict. Point 3 remains [MANUAL_REVIEW](#) for every node either way. A non-match means the page body does not carry the approved string — which may mean the page is wrong, that the name is shown in an image or the `<title>` rather than in text, or that the official list is out of date. Deciding which is the reviewer's job; the tool's job is to say precisely what it did and did not find. Where nothing matched, the evidence also reports how many times the phrase the approved names share ([EOSC Node](#)) does appear, so you know where to look.

If a name you expect does not match, read [full_text](#) in that node's evidence file before assuming the page is at fault.

Adding a node

Three files change together, and the test suite refuses to let them disagree. Nothing is generated or auto-discovered: every value below is typed in by a human, on purpose, so that the report can say where it came from.

1. [nodes.yaml](#) — add the entry. Put the Node Landing Page in [url](#), taken from the **EOSC EU Node Contributors Dashboard** field "Website address" (section 1.2, field 6), not from a search engine and not from the node's general homepage. That registered URL is what the checklist defines as the Node Landing Page, so if what you were sent differs from what is registered, the registered value wins.

Then look up [eosc_page](#) on the live index at [eosc.eu/building-the-eosc-federation/](#) and copy the slug. Do not guess it from the node's name: the slugs are not formed consistently — Czechia is [eosc-node-czechia](#), but the Digital Twin of the Ocean is [eosc-node-digital-twin-of-the-ocean](#) while its node id here is [eosc-dto](#). A wrong slug makes point 4 look for a page that does not exist, and the node fails a requirement it may well satisfy.

2. [checklist/approved-names.txt](#) — add the node's approved name, exactly as the Tripartite list writes it, including the [EOSC Node |](#) prefix. This file is the official list as circulated; do not edit the wording to match what a page happens to say. If the official list has not yet been updated for the new node, leave this file alone and say so — a name absent from the official list is a real finding, and inventing an entry would hide it.

3. [checklist/approved-names-scoped.txt](#) — add the same name, prefixed with the node id you chose in step 1: [cern: EOSC Node | CERN](#). The names must match the official file character for character; only the id prefix is added here.

Then run the suite **before** running the checker:

```
uv run pytest -q
```

The guards in [tests/test_nodes.py](#) and [tests/test_checklist.py](#) will tell you exactly which file you missed — a node with no scoped name, an id that is not a usable filename, a URL that is not [https](#), an [eosc_page](#) that is not on [eosc.eu](#) or that points at the federation index rather than the node's own entry, two nodes sharing a landing page, or a name present in one list and not the other. These fail in milliseconds, before anything touches the network.

Only then collect the new node. Fetch just the node you added, rather than re-running the whole federation, and into a scratch directory first:

```
uv run basic-check collect --only <new-id> --results /tmp/trial
uv run basic-check assess --only <new-id> --results /tmp/trial
```

How to move that evidence into the published run without fetching the page a second time is [example 1](#) in [section 11](#).

`assess` reads evidence per node from disk, so a node configured but not yet collected is **not** silently skipped: the run prints `N node(s) have no evidence and were NOT assessed`, names them, and **exits 2** while still writing the report for the rest. That exit code is the signal that the published report is incomplete — do not commit a report produced by a run that exited 2 unless you intend to publish a partial one and say so.

If you maintain your own copy elsewhere (GitLab, or a fork)

The three files above are the whole of it, and none of them is specific to GitHub. `nodes.yaml` and both `checklist/*.txt` files are ordinary version-controlled text at the repository root; edit them on whatever host you use, in a branch or straight on the default branch, and the checker behaves identically. There is no registry, no database, no service to notify, and no value cached anywhere else in the repository.

Two host-specific things do **not** carry over:

- `.github/workflows/` is GitHub Actions only. On GitLab you would need a `.gitlab-ci.yml` expressing the same two jobs — install with `uv sync --locked`, run `pytest`, and optionally run the checker on manual trigger only. Keep the manual-trigger restraint: a scheduled compliance job means fetching other organisations' production websites on a timer.
 - The committed `results/` directory is this repository's published output. A separate copy will produce its own, and the two will diverge as the sites change. Decide which copy is authoritative before both are quoted in a meeting.
-

4. Running it

```
uv run basic-check run
```

That is the whole thing: fetch every configured node, apply the checklist, write the reports. It takes a few minutes, most of it spent deliberately waiting between requests.

The commands

```
uv run basic-check points      # the checklist, and what is decidable
uv run basic-check collect     # fetch -> results/evidence/
uv run basic-check assess      # evidence -> results/
uv run basic-check run         # collect, then assess
uv run basic-check show eudat   # one node's results in the terminal
```

Listing the configuration

Six flags on `basic-check` itself, used without a command, show what is configured:

```
uv run basic-check --list-nodes      # node ids, one per line
uv run basic-check --list-nodes-ids  # the same, under a second name
uv run basic-check --list-nlps       # id and Node Landing Page URL
uv run basic-check --list-approved-names # id and approved name
uv run basic-check --print-config     # all three in one table
uv run basic-check --show-node bbmri-eric # one node's row of that table
```

`--print-config` prints one row per node, in `nodes.yaml` order, as plain aligned columns (shortened here):

Node id	Node Landing Page URL	Approved name
bbmri-eric	https://www.bbmri-eric.eu/eosc-node-bbmri-eric/	EOSC Node BBMRI-ERIC
cern	https://eosc-auth.cern.ch/login	EOSC Node CERN
...	...	...
eosc-sk	https://eosc.sk/	EOSC Node Slovakia

followed by two lines naming the files: `* 13 node(s). Ids and URLs from nodes.yaml; approved names from checklist/approved-names-scoped.txt.` and `"assess and run use checklist/approved-names.txt by default: the same names, unscoped, so each counts for every node."`

`--show-node` takes one node, by id or by name in any case as `--skip` does (`eosc-it`, `Italy`, `EOSC Node Italy`), and prints that node's row of the table under the same header:

Node id	Node Landing Page URL	Approved name
bbmri-eric	https://www.bbmri-eric.eu/eosc-node-bbmri-eric/	EOSC Node BBMRI-ERIC

A value that matches no node, or more than one, is an error that lists the ids. `--list-nodes-ids` is another name for `--list-nodes`, with the same output.

The approved name next to each id comes from `checklist/approved-names-scoped.txt`, the copy of the official list that ties each name to a node. `assess` and `run` use `checklist/approved-names.txt` by default, the same names unscoped, so there each name counts for every node; the table says so under the rows. A node with no name in the scoped file shows `(none)`. The flags can be combined (`--list-nodes --list-nlps`), print plain columns that are never wrapped, so they can be piped to `grep` or `cut`, and are refused together with a command. They read local files only and contact no node.

Use it before and after editing `nodes.yaml` or the name lists, to see what the next run will check.

Useful options

Not every option is accepted by every command, and passing one to the wrong command is an error rather than a no-op. Each table below names the commands that accept the option.

Fetching — `collect` and `run` only, because only these two touch the network:

Option	Effect
<code>--depth N</code>	How far to follow links. See section 5.
<code>--delay S</code>	Seconds between hosts (default 2.0). Raise it to be gentler.
<code>--fetch-budget N</code>	Depth 2 only: hard ceiling on second-hop requests for the whole run (default 60).
<code>--max-children N</code>	Cap on followed pages per node at depth 1 (default 8). collect only — run rejects it , so use <code>collect</code> then <code>assess</code> if you need it.

Choosing what to check — accepted by `collect`, `assess` and `run`:

Option	Effect
<code>--only a,b</code>	Restrict which nodes are fetched . The report still covers all of them — see below.
<code>--skip a,b</code>	Leave these nodes out entirely: not fetched, not assessed, not in the report, which says so. Id or name, comma-separated or repeated — see below.
<code>--url https://...</code>	Check any page without editing <code>nodes.yaml</code> . Repeatable. With <code>--node</code> , the alternative page for that node.
<code>--node id</code>	Check one configured node at the alternative landing page given with <code>--url</code> , keeping its id, name and <code>eosc_page</code> . See below.
<code>--eosc-page URL</code>	With a single <code>--url</code> : that node's own <code>eosc.eu</code> page, so a point 4 failure can name the exact URL that is missing. With <code>--node</code> , replaces that node's configured one.
<code>--nodes path</code>	Use a different node list.

Assessment and output:

Option	Where	Effect
<code>--approved-names path</code>	assess, run	Use this name list instead of the committed <code>checklist/approved-names.txt</code> . One name per line, ideally as <code>node-id: Name</code> . A path that does not exist is an error, not a fall back to the default. See section 3.
<code>--no-approved-names</code>	assess, run	Use no name list at all, not even the committed default. Point 3's name requirement is then not assessed.
<code>--strict-separators</code>	assess, run	Match the separator glyph in each approved name literally. By default a page writing a hyphen satisfies a list writing a pipe; with this flag it does not. See section 3.
<code>--checklist path</code>	assess, run, points	Use a different checklist version.
<code>--run LABEL</code>	assess, run	Label stored with the run, for telling one report from another.
<code>--results path</code>	all but points	Write somewhere other than <code>results/</code> .
<code>--one-off</code>	show only	Read <code>results/one-off/</code> instead of the federation run.

`--only` narrows the fetch, not the report

Earlier versions of this tool had a real defect here, and this guide documented it as a warning: `assess --only data-terra` replaced the full matrix in `results/` with a single row, and the resulting `index.html` looked like a federation report that covered one node. A short table is indistinguishable from a complete one.

That is fixed. `--only` now narrows **which sites are contacted**, and the report written to `results/` still covers every node in `nodes.yaml`:

```
uv run basic-check assess --only data-terra
# One node re-assessed from fresh evidence; results.json still has every node.
```

The rows you did not select are reused from the evidence already on disk, so the report is no longer uniformly fresh. It says so itself: a **Mixed freshness** banner names the nodes that were re-fetched, states how many were not, and gives the capture date of the reused evidence. `results.json` records the same thing in a top-level `selection` list, which is `[]` on a full run.

An explicit `--results DIR` keeps the narrowing behaviour, because there nothing shared is at risk and a one-node scratch table is usually the point. Note that `--results DIR` relocates the whole workspace, not just the output: evidence is read from `DIR/evidence` as well as written there. Point it at an empty directory and you get an empty report and exit code 2, not a one-node one. Give it evidence first:

```
mkdir -p /tmp/scratch && cp -r results/evidence /tmp/scratch/evidence
uv run basic-check assess --results /tmp/scratch --only data-terra
# A deliberate one-node report, written somewhere of your choosing.
```

Or collect straight into it, which re-fetches that node rather than reusing anything:

```
uv run basic-check collect --results /tmp/scratch --only data-terra
uv run basic-check assess --results /tmp/scratch --only data-terra
```

A bare `assess` still rebuilds every row from evidence on disk, so it remains the way to produce a report with a single capture date:

```
uv run basic-check assess
```

`--skip` leaves nodes out, and says so

`--skip` is for the opposite case: run on every configured node **except** some. It takes node ids or names, case-insensitively, either comma-separated or by repeating the option. Short names work too, so `Italy` names "EOSC Node Italy":

```
uv run basic-check run --skip Italy,Slovakia
uv run basic-check run --skip Italy --skip Slovakia    # the same
uv run basic-check assess --skip "CERN, Czechia" --skip eossc-it --skip eossc-sk
```

A skipped node is left out of the whole run. No request is sent to its site, it is not assessed, and it has no row in the report. Because it is left out on purpose, it does not count as missing evidence, so `assess` does not exit 2 for it. The reviewed report of 24 September was built this way, with `--skip Italy`, because `eosc.it` had no address record that day.

Unlike `--only`, `--skip` does shorten the report, including the one in `results/`. A shorter table must never look complete, so every report says what was left out: a **Skipped by request** banner in `results.md` and `index.html`, and a top-level `skipped` list in `results.json` (`[]` when nothing was skipped). Like any `assess` without `--results`, it rewrites `results/`, which stays under human review. For a trial run, add `--results /tmp/scratch`.

The option rejects values rather than guessing. If a value matches no node, the error lists the valid ids, because a typo that skipped nothing would fetch the very site you meant to leave alone. A value that names more than one node is an error too. So is skipping every node, and so is combining `--skip` with `--url`. With `--only`, the skip is applied after the selection: `--only egi,eudat --skip egi` fetches EUDAT alone.

Checking a page not in `nodes.yaml`

```
uv run basic-check run --url https://example.org/our-node-page
```

Ad hoc checks write to `results/one-off/` so a published run is never overwritten, and the report carries a banner saying it is not a federation run. That banner matters: these reports circulate before a production decision, and a stray single-page file that looks official is a genuine hazard.

Updating one node's row: `--update-results-for-node`

After a node's URL changes, or when one node's page has been fixed, there is no need to re-check the whole federation. `--update-results-for-node NODE`, on `run` and `assess`, re-checks that node alone and replaces only its row. The stored `results.json` in the results directory is the memory: every other row is copied from it as it is, neither re-fetched nor re-judged.

Command	What it does
<code>run --update-results-for-node NODE</code>	Fetches that one page again, into a scratch directory. Only if the page loaded does it replace the node's evidence and screenshot, and its row. <code>--accept-error</code> merges a failed fetch too.
<code>assess --update-results-for-node NODE</code>	Judges the node's evidence already in <code>evidence/</code> again. Offline.

Either way, `results.md`, `index.html` and `results.csv` are rebuilt from the stored rows, and the command prints each point's old and new verdict. `NODE` is an id or a name, in any case. A node the stored run skipped is inserted where `nodes.yaml` puts it and removed from the skipped list. The run id and time in the header stay those of the original run; `results.json` gains `updated_at` and a `row_updates` list with the node's previous URL, capture time and verdicts, and both reports carry an **Updated rows** banner saying which row is newer and that the others were not re-judged.

It refuses, and changes nothing, when:

- the results directory holds no `results.json`, or holds an alternative-URL trial (`--node`);
- the checklist is not the one the stored results were assessed with, or the approved-name options (`--approved-names`, `--no-approved-names`, `--strict-separators`) differ from theirs. One row judged by other rules would not be comparable with the rest; run the whole federation again instead;
- it is combined with `--only`, `--skip`, `--url`, `--node`, `--eosc-page` or `--run`.

The default directory is `results/`, the published run. Try it on a copy first:

```
cp -r results /tmp/copy
uv run basic-check run --update-results-for-node bbmri-eric --results /tmp/copy
uv run basic-check show bbmri-eric --results /tmp/copy
```

Checking a node at an alternative URL

To run every test against one configured node at a landing page other than the one in `nodes.yaml`, give the node with `--node` and the page with `--url`:

```
uv run basic-check run --node bbmri-eric --url https://alt.example.org/eosc-node-bbmri-eric/
uv run basic-check show bbmri-eric --one-off
```

The node keeps its id, name and `eosc_page`, so its scoped approved name applies to point 3 and point 4 looks for its own `eosc.eu` page; only the page fetched changes, and `nodes.yaml` is not edited. `--node` takes an id or a name in any case, as `--skip` does.

Rule	Why
Exactly one <code>--url</code>	It is the alternative page for that one node. Without <code>--url</code> , use <code>--only</code> to check the node at its configured page.
Not with <code>--only</code> or <code>--skip</code>	The run is already one node; either would be ignored or contradict it.
Writes to <code>results/one-off/</code> , or another <code>--results</code> folder	<code>--results results/</code> is refused: the evidence file is named after the node id and would replace the reviewed evidence.
<code>--eosc-page</code> replaces the node's <code>eosc.eu</code> page	For a trial of a changed registration as well as a changed page.

The reports are titled **Alternative URL check (not a federation run)** and carry an **Alternative URL, not a federation run** banner naming both addresses; `results.json` records the pair under `alternative_url`, each node record keeps `configured_url`, and `show` prints both. If the same evidence is later assessed without `--node`, the usual **Evidence from a different URL** warning appears, so it cannot pass as the configured page.

This is the way to try a node's new address before changing `nodes.yaml` (example 2 in section 11).

5. How deep to go: `--depth`

<code>--depth</code>	What is fetched	Requests, 9 nodes, 21 Sep 2026	Requests, 12 nodes, 24 Sep 2026	Requests, 13 nodes, 26 Sep 2026
0	The landing page only.	9	12	13
1 (default)	The landing page, plus links from it that could settle a checklist point — policy, contact, about. At most 8 per node, 2 per point.	31 (9 + 22)	44 (12 + 32)	44 (13 + 31), the published run

<code>--depth</code>	What is fetched	Requests, 9 nodes, 21 Sep 2026	Requests, 12 nodes, 24 Sep 2026	Requests, 13 nodes, 26 Sep 2026
2	The above, plus one further hop: a policy <i>index</i> that links on to the actual policy, for instance. At most 2 per fetched page and 1 per point, under a run-wide budget.	45 (9 + 22 + 14)	not run	not run

[robots.txt](#) requests are not counted in these figures. The tool fetches one per host before touching it.

Links are followed only where the target could settle a point, so a policy link is verified rather than taken on the strength of its label.

```
uv run basic-check run --depth 2 # bounded by the default budget
uv run basic-check run --depth 2 --fetch-budget 20 # stricter ceiling
```

A depth-2 run reports both depths. You get two summary tables — *Results at depth 1* and *Results at depth 2* — two explanatory sections, a list of every cell whose verdict differs, and a table of exactly which second-hop pages were fetched and what each was followed for.

Both tables come from the **same capture**: the shallow view is the deep evidence with the second-hop pages set aside, not a second run. Nothing is fetched twice to produce the comparison, so any difference between the two tables is the extra hop and not the page changing in between.

Is depth 2 worth it? On this federation, so far, no. Two depth-2 runs on 21 September (18 and 14 second-hop requests) each changed **zero** verdicts. That is a finding, not a disappointment: the points still marked review turn on a judgement ("clearly state") or quantify over things no crawl enumerates ("all research resources offered by the Node"), and no amount of fetching settles either kind. Run it yourself before assuming the same holds for another federation or a later date.

Be considerate with other people's servers

These are production sites that did not ask to be tested. The defaults are deliberately restrained: one request per node every two seconds, a per-node cap on followed links, and a single run-wide budget for second-hop requests rather than a per-node one.

This is not theoretical. [geant.org](#) served [HTTP 200](#) at 12:43 UTC on 21 September, [HTTP 403](#) at 13:07 the same afternoon after several runs, redirecting to a Cloudflare bot-protection challenge, and [HTTP 200](#) again at 18:39 from a GitHub runner — a different address that had made no requests that day. It happened again on the 25 September trial, at 04:42 UTC: GÉANT answered HTTP 403 with the same challenge, and five of its cells that were decided on 24 September moved to review. **Repeated automated runs are visible to the sites you are checking**, and the block tracks the requesting address rather than the tool's identity. Run the full suite when you need a result, not in a loop, and use [assess](#), which needs neither browser nor network, when you are iterating on the report itself.

6. Reading the output

results/ after a run:

File	What it is
index.html	The deliverable. The matrix, per-node detail, and the full checklist text. Open it in a browser.
results.md	The same content as Markdown, for reading on GitHub or pasting into a document.
results.csv	One row per node-point, for a spreadsheet or pivot table.
results.json	The full structured run, including every piece of evidence. Use this for any further analysis.
checklist-v3.1.html	The checklist points in full, with the reasoning behind each verdict.
evidence/<id>.json	The raw capture for one node: links, images, controls, text, headers.
evidence/screenshots/<id>.png	What the page looked like when it was fetched — the visible viewport only , so footers and anything else below the fold are not in the image. Do not read a missing element here as absent from the page; check links and full_text in the evidence JSON.

The four verdicts

	Meaning
 PASS	Satisfied, and the tool can show why.
 FAIL	Violated, and the tool can show why.
 review	A human must decide. The tool has gathered the evidence and refuses to guess.
 ERROR	Could not be assessed at all.

ERROR is not a finding against the node. It means the tool has no page to judge: the host did not resolve, the request failed, or [robots.txt](#) asked crawlers to stay away and the tool complied. The reason is in each cell's message, for example [The page could not be fetched: not fetched: robots.txt disallows it](#), and [robots_note](#) in the evidence file records what [robots.txt](#) said. In the 24 September run only Italy would have been **ERROR**, and it was skipped instead; the published run of 26 September has no **ERROR**. On the 25 September trial, BBMRI-ERIC's new [dev3](#). address was **ERROR** on all ten points because its [robots.txt](#) disallows every path.

A FAIL is an accusation, so the tool is careful about making one. Absence is never concluded from a page that did not render: if a capture yields no links at all, or fewer than ten DOM elements, or under 500 characters of text, the tool returns review rather than FAIL. The point is that "we found no contact link" and "the page did not load for us" must not produce the same verdict.

The same restraint applies to being blocked. GÉANT's [HTTP 403](#) moved four cells from decided verdicts to review — the tool does not convert bot protection into a compliance failure. A later run from an unthrottled address reached the page and settled all four, which is the restraint paying off: had the block been recorded

as failure, that report would have contained four wrong verdicts instead of four honest abstentions. The block came back in the published run of 26 September, and the same four cells are review there.

Reproducing a verdict by hand

Every verdict is traceable to evidence on disk:

```
uv run basic-check show geant # verdicts and evidence in the terminal
python3 -m json.tool results/evidence/geant.json | less # the capture, personal data masked
```

results.json carries the same evidence strings the report displays, so you can check any claim without re-fetching anything.

To work out why a particular cell came out the way it did, the decision procedure for every point — every branch, in evaluation order, with the verdict it yields — is tabulated in section 5 of the analysis workflow. Section 6 of that document lists the cases where the tool is known to be wrong, which is the first place to look when a verdict surprises you.

Personal data is masked

Contact pages sometimes name individuals, such as a media officer with a direct line. The checklist needs none of that: point 6 asks for a way to reach the node helpdesk, which is a role. So src/basic_check/privacy.py masks personal data at the two places anything is written to disk. When evidence is collected, the evidence file is masked before it is saved. When the reports are written, the run is masked again before any of the four formats is rendered, so a report rebuilt from evidence captured before masking existed is masked too.

Found on the page	Written as
jane.doe@example.org, jane.doe [at] example.org	XXXXXX@example.org, XXXXX [at] example.org
support@egi.eu, it@helpdesk.bbmri-eric.eu, geant@geant.org	unchanged: role and organisation mailboxes
+31 20 123 4567, tel:0043316349917	+31 XXXXXX, tel:+43 XXXXXX
tel:+31(0)20%20123456 (a space written %20 inside a link)	tel:+31 XXXXXX
Tel.: 06 1234 5678, tel. (09) 123 4567 (no prefix, but labelled)	Tel.: XXXXXX, tel. XXXXXX
Mgr. Jana Nováková, novakova@...	Mgr. XXXXX XXXXX, XXXXX@...

An address is kept only when it is recognisably a role, such as support, info, helpdesk, csirt or any mailbox under a helpdesk. host. Every other address is masked. A name is masked only near a masked address, and never inside a URL or as an all-capitals acronym. Office switchboards are masked like any other number, because nothing in a number says who answers it. Masking changes no verdict, because every check works from the parts that are kept. Screenshots are not masked; they show the visible part of the landing page only, not the contact pages where these details appeared.

The published run is masked. Run web-5 (26 September 2026) was collected with masking in place, and a search of all its evidence and reports found no personal address or phone number. Before it, commit ada1b4a (25 September 2026) masked the 24 September run's results/evidence/ and rebuilt the reports offline. All 120 verdicts and the 44 / 6 / 70 tally were unchanged. A few character counts in the report dropped slightly, by the length of the masked text (BBMRI-ERIC's page from 7100 to 7094 characters). The first masking pass had let two phone formats through, the %20 and the bracketed area code in the table above. Both were found by searching every published file for phone-shaped text, fixed, and added to the

tests. `tests/test_privacy.py` now also checks the committed evidence itself, and fails if any file in `results/evidence/` contains a personal address or a phone number. The earlier, unmasked versions remain in the git history. Removing them would mean rewriting the history of `main`, which has not been done.

To mask evidence collected before masking existed, for instance a trial run in a scratch directory, re-write it with the same function and rebuild:

```
uv run python -c "
import json, pathlib
from basic_check.privacy import mask_data
for p in pathlib.Path('/tmp/trial/evidence').glob('*.json'):
    p.write_text(json.dumps(mask_data(json.loads(p.read_text())), indent=2, ensure_ascii=False))
"
uv run basic-check assess --results /tmp/trial --run <run id>
```

Rebuild with the `nodes.yaml` the run was collected with (`--nodes`), or a URL changed since then is shown against evidence taken from the old one (section 3, "Changing a node's URL").

7. The test suite

```
uv run pytest -q           # 426 tests, offline, a few seconds
uv run pytest -v           # names of every test
uv run pytest tests/test_checks.py # one file
uv run pytest -k depth     # anything about depth
uv run ruff check src tests # lint
```

File	Tests	Covers
<code>test_checklist.py</code>	15	The transcription matches the source document, including its SHA-256 for every committed revision, the scoped name list agrees with the official one, and switching <code>DEFAULT_CHECKLIST</code> carries the help texts with it.
<code>test_checks.py</code>	75	The verdict logic, point by point, including the render gate, the EOSC-asset token rule, that a point 3 summary never denies having a name list it was given, and that a link merely labelled "support" does not settle point 6.
<code>test_cli.py</code>	113	Command wiring, options, ad hoc <code>--url</code> isolation, that <code>--only</code> does not shrink the published report, that <code>--skip</code> leaves nodes out and says so, that <code>--help</code> stays complete, and that evidence from another URL than the configured one is flagged.

File	Tests	Covers
test_crawl.py	35	Link selection, including policy words found only in a hyphenated address, host containment, depth-2 budget, the depth-1 view.
test_names.py	64	Parsing, scoping, word boundaries, separator flexibility and its strict counterpart, and the recorded digest.
test_nodes.py	14	nodes.yaml itself, and that the default copy shipped for <code>--restore-default-config</code> is the committed file: every node declares every field, ids are unique and usable as filenames, URLs are absolute https , no two nodes share an eosc_page , no eosc_page is the federation index, and every node has a scoped approved name.
test_privacy.py	51	Personal-data masking: personal addresses and phone numbers masked, role mailboxes kept, the name beside an address masked but not a title or an acronym, that both the evidence files and every report format are written masked, and that the committed evidence stays masked.
test_report.py	36	Matrix rendering, the dual-depth tables, the mixed-freshness banner, the name-list provenance, and input that would break a table or a list — a or a newline in a node name, an evidence line, a followed-link reason or a point title.
test_update_row.py	22	<code>--update-results-for-node</code> , on a copy of the committed results/ : only the named row changes, the reports gain one Updated rows line, run collects exactly one node, a failed fetch changes nothing unless <code>--accept-error</code> , a skipped node is inserted in nodes.yaml order, and every refusal.

The suite makes no network requests and needs no browser, which is why CI runs it without downloading Chromium.

A fuller treatment — why the suite takes a few seconds rather than a fraction of one, the testing criteria behind it, and the defects that each earned a regression test — is in the [test suite and configuration overview](#).

Test names are sentences, because a failing test should tell you what broke without opening the file:

```
test_a_consent_banner_does_not_excuse_a_missing_link
test_the_two_tables_show_their_own_verdicts
test_a_pipe_in_a_url_cannot_break_the_table
```

A note on what a passing suite means

Green means nothing objected. It does not mean everything was verified.

When adding a check, make the test fail first, then implement — and when changing a guard, mutate it deliberately to confirm the test actually bites. A test that cannot fail is worse than no test, because it produces confidence without evidence. Both dual-depth guards in [test_report.py](#) were confirmed this way: rendering the deep results in the shallow table, and forcing the stray heading back, each fail exactly one test.

8. Continuous integration

[.github/workflows/tests.yml](#) runs on every push and pull request, and can also be started by hand: `uv sync --locked`, `ruff check`, then `pytest` with a JUnit report published to the run summary and uploaded as an artifact. No browser is installed, because the suite does not need one.

Both workflows use `--locked` rather than a plain `uv sync`. The difference matters: plain `uv sync` quietly updates the lockfile when it disagrees with `pyproject.toml` and carries on, so a dependency change could go green in CI while the versions tested were not the versions committed. `--locked` installs exactly what `uv.lock` pins and fails otherwise. If CI ever stops at the install step saying the lockfile needs updating, that is the guard working — run `uv lock` and commit the result.

[.github/workflows/compliance.yml](#) is **manual dispatch only** — deliberately. A scheduled compliance run would mean fetching every node's production website on a timer, which is exactly the behaviour that got GÉANT's bot protection to start refusing requests. The workflow reads [nodes.yml](#) from the branch it runs on, so on `main` today it fetches BBMRI-ERIC's public [www](#). address, the one the published run was collected from.

It takes three inputs: `only` (comma-separated node ids, empty for every node), `depth` (1, 0 or 2) and `delay` (seconds between nodes). `depth: 2` was added after a web run could not reproduce the published depth-2 report — the input had offered only 1 and 0, so the workflow could not produce the report the repository publishes. Choose it when you want that like-for-like comparison, and note it means roughly fourteen further requests to other people's sites.

The workflow has `permissions: contents: read` and does **not** commit its output. The reports are uploaded as the `compliance-results` artifact, and republishing is a deliberate local step: download the artifact, review it, and commit. A workflow that could rewrite the published verdicts unattended would remove the review that makes them worth publishing.

Its run summary publishes **counts only** — how many nodes, and the tally of verdicts — labelled as unreviewed automated output, with the per-node detail left to the [compliance-results](#) artifact. On a public repository the run summary is world-readable, and while the committed report is public anyway, a table of FAILs against named organisations generated automatically under the repository's name reads as a finding rather than as a draft. The summary also declines to report anything unless the check step succeeded, because [results.json](#) is committed: a failed run would otherwise have read the previous reviewed run's counts out of a fresh checkout and published them under a new run number.

9. Troubleshooting

Symptom	Cause and fix
Executable doesn't exist ... playwright install	Chromium is not installed. Run uv run playwright install chromium --with-deps . Not needed for assess , show , points or pytest .
The report has fewer nodes than you expected	You passed --results DIR together with --only , which narrows both deliberately. Without an explicit --results , --only narrows the fetch alone and the report keeps every node. Section 4.
The report carries a Skipped by request banner, or has fewer rows than nodes.yaml	Expected after --skip : those nodes were left out on purpose, and the banner and results.json name them. Run without --skip for every node. Section 4.
The report carries a Mixed freshness banner	Expected after --only : the nodes you did not select were reused from evidence on disk. Run a bare uv run basic-check collect then assess for a report with one capture date. Section 4.
Every point for a node is ERROR , with the message not fetched: robots.txt disallows it	The site's robots.txt excludes the tool, and it complied. Common on staging hosts: BBMRI-ERIC's dev3 . address does it. Ask the node for the public address or an exception; do not work around it. Section 6.
assess warns that nodes were assessed from evidence collected at a different URL, and the report carries an Evidence from a different URL banner	A node's url in nodes.yaml changed after its evidence was collected, so the row shows the new address above the old page's verdicts. Collect that node again, or rebuild with the old node list via --nodes . Section 3, "Changing a node's URL", and example 2 in section 11.
test_the_committed_evidence_publishes_no_personal_address_or_phone fails	Evidence that is not masked has reached results/evidence/ , usually copied in from a run made before masking existed. Mask it with mask_data and rebuild (section 6). Do not commit until the test passes.
A node shows HTTP 403 and everything moved to review	Bot protection, not an access policy. Check final_url in the evidence for a __cf_chl_rt_tk parameter. Wait, run less often, or verify that node by hand in a browser.

Symptom	Cause and fix
Everything is review for one node	The capture probably did not render. Look at evidence/screenshots/<id>.png — that is exactly what the render gate is protecting you from.
<code>test_checklist.py</code> fails on a hash	The checklist PDF changed. That is the test doing its job: transcribe the new version into a new YAML file rather than adjusting the hash.
A run takes far longer than expected	<code>--delay</code> defaults to 2.0s between hosts and slow nodes are waited on. This is intentional.
Point 3 reports <code>NONE of the ... approved name(s) ... appear</code> for a node you know is named correctly	Expected for many nodes: in the published run the committed list matches 4 of 13. The <code><title></code> is not searched, and a name is not matched inside a longer word — though the separator between words is flexible. Read full_text in that node's evidence file. Section 3.
Point 3 says <code>no approved name was supplied for this node</code>	Your list is scoped and has no <code>node-id:</code> line for that node. Add one, or use a bare name to cover every node. Section 3.
Point 4 says <code>PASS</code> but you cannot find the link	Look in the footer. The check reads the DOM, not the visible area, and several nodes put the eosc.eu link in a legal/navigation column at the very bottom. The exact URL is quoted in the evidence — search the page for it rather than scanning by eye. Note also that the target is never requested, so the check cannot tell you the page still exists.
Point 3 says <code>No approved-name list was used</code>	You passed <code>--no-approved-names</code> , or checklist/approved-names.txt is missing from your checkout (the run warns on stderr when it is). Section 3.
The report says the name list came from <code>a list supplied for this run</code> when you expected the official one	You passed <code>--approved-names</code> . Drop the flag to use the committed default. Section 3.
Point 3 says a name <code>was not looked for</code>	The page returned no body — a 403 or a failed render. Fix the collection first; the name says nothing until there is a page to read. Retrying later, or from a different address, is often enough.
A match is reported <code>from the unscoped list</code> and you want a firmer claim	Prefix each name with its node id so it is only matched against that node. Section 3.

10. Extending it

- **Add a node:** three files together — `nodes.yaml`, [checklist/approved-names.txt](#) and [checklist/approved-names-scoped.txt](#) — then `pytest`, then `collect --only <id>`. Full procedure and failure modes in [Adding a node](#) in section 3; the commands through to publication in example 1, section 11.
- **Change a node's URL:** the `url` line in `nodes.yaml` only. Example 2, section 11.

- **Add a checklist version:** new YAML beside `v3.1.yaml`, with `source_file` and `source_sha256`; do not edit an existing version in place. Then change `DEFAULT_CHECKLIST` in `src/basic_check/cli.py`. Example 3, section 11.
- **Add a check:** implement in `src/basic_check/checks.py`, write the failing test first, and prefer returning `MANUAL_REVIEW` with good evidence over a confident guess. Document its branches in `ANALYSIS-WORKFLOW.md` — a check whose decision procedure is not written down cannot be reviewed.
- **Change the report:** `src/basic_check/report.py` renders HTML, Markdown and CSV from one run dict. `tests/test_report.py` covers the matrix; add to it, because a rendering bug is silent.

Text that reaches the Markdown report is escaped at the boundary. Two helpers in `report.py` do it, and new rendering code should use them rather than interpolating a value directly:

Helper	Use for	What it does
<code>_md_cell(value)</code>	Anything placed in a table cell	Escapes the pipe character so it cannot open a column, and flattens newlines
<code>_md_text(value)</code>	Inline text — list items, headings, bold runs	Flattens newlines so the element does not end early

Both collapse runs of whitespace, which leaves ordinary text untouched. A pipe in a **URL** is the exception: inside an autolink a backslash is not an escape, so pipes there are percent-encoded to `%7C` instead.

The reason this is handled centrally is that a rendering bug of this kind is silent — the report is still valid Markdown, just with the columns shifted or a sentence promoted out of its list item, and nothing raises. Node names come from `nodes.yaml`, evidence strings from fetched pages, and point titles from a checklist transcription, so none of it is under this module's control.

11. Worked examples

Three step-by-step procedures, one for each change that comes up most often: adding a node, changing a node's URL, and adopting a new revision of the checklist. Each lists which files to edit, the exact commands, and how to put the result into the published `results/` directory.

Four rules apply to all three:

- **Try it in a scratch folder first.** `--results /tmp/trial` writes nowhere near `results/`, so nothing published changes until you decide it should.
- **Run `uv run pytest -q` after every edit.** It needs no network, takes a few seconds, and names the file you got wrong.
- **Contact as few sites as possible.** Only `collect` and `run` send requests. `assess` works on saved evidence, so it can be repeated as often as you like.
- **A new published run needs a human review before it is committed.** The tool's verdicts are not a compliance statement (section 6). Write the review down, as `results/REVIEW-2026-09-24.md` does for the 24 September run, and commit only then.

The run labels below (`live-2026-10-01-...`) are examples. Use the date of your own run.

Example 1 — adding a node and publishing it

The example adds a node with the placeholder id `eosc-example`. Replace every value with the real ones.

Step 1: edit three files. All three are at the repository root or in `checklist/`, and all are edited by hand.

nodes.yaml: add the entry. `url` is the landing page as registered in the EOSC EU Node Contributors Dashboard ("Website address", section 1.2, field 6). Copy `eosc_page` from the [federation index](#); do not guess it.

```
- id: eosc-example
  name: EOSC Node Example
  url: https://eosc-node.example.org/
  eosc_page: https://eosc.eu/building-the-eosc-federation/eosc-node-example/
```

checklist/approved-names.txt: add the approved name exactly as the Tripartite list writes it:

```
EOSC Node | Example
```

checklist/approved-names-scoped.txt: add the same name, with the node id in front:

```
eosc-example: EOSC Node | Example
```

If the Tripartite list does not include the node yet, do not make a name up. Leave both name files alone. `test_nodes.py::test_every_node_has_a_scoped_approved_name` then fails and names the node, which is intended: the node is not ready to be published. In the meantime, check it in a scratch folder with `--url https://eosc-node.example.org/` (see [Checking a page not in nodes.yaml](#), section 4), which needs no configuration.

Step 2: check the configuration, offline.

```
uv run pytest -q
```

A missing scoped name, an id that is not a valid filename, a URL that is not [https](#), an `eosc_page` outside [eosc.eu](#), or a landing page already used by another node each fail with a message naming the problem.

Step 3: trial run for the new node only. This contacts one site.

```
uv run basic-check collect --only eosc-example --results /tmp/trial
uv run basic-check assess --only eosc-example --results /tmp/trial
uv run basic-check show eosc-example --results /tmp/trial
```

Open `/tmp/trial/index.html` or `/tmp/trial/results.md` and read the new row. If the page returned **ERROR** (robots.txt, a 403, a DNS failure), stop there. An **ERROR** row is not a result worth publishing.

Step 4: put the node into the published run. There are two ways.

A — add the one node to the current run (one site contacted, and already done). Reuse the evidence from step 3 rather than fetching the page again:

```
cp /tmp/trial/evidence/eosc-example.json results/evidence/
cp /tmp/trial/evidence/screenshots/eosc-example.png results/evidence/screenshots/
uv run basic-check assess --only eosc-example --skip Italy \
  --run live-2026-10-01-plus-eosc-example
```

The report covers every node. A **Mixed freshness** banner says that only `eosc-example` is new and that the other rows are reused from the earlier capture, with its date. `--skip Italy` keeps the same scope as the published run. Italy's URL changed on 26 September to <https://eoscnode-it.d4science.org/>; drop `--skip Italy` once a trial collection (`collect --only eosc-it --results /tmp/trial`) shows the page loads.

B — a complete new run (every node contacted once). Use this when the existing evidence is old enough that a single timestamp is worth more than leaving the other sites alone:

```
uv run basic-check run --skip Italy --results /tmp/new-run --run live-2026-10-01
# review /tmp/new-run/results.md and index.html, then:
rm -rf results/evidence
cp -r /tmp/new-run/evidence results/evidence
uv run basic-check assess --skip Italy --run live-2026-10-01
```

The last command re-renders the reports in `results/` from the evidence you reviewed. It makes no requests.

Either way, `assess` must end with **exit code 0**. Exit 2 means a configured node has no evidence, so the table is incomplete.

If `assess` warns **"Evidence from a different URL"**. It means a node's URL in `nodes.yaml` is not the page its evidence was collected from. That is the case for BBMRI-ERIC from 25 to 26 September 2026, when `nodes.yaml` had the `dev3`. address but the published evidence was from the `www`. one. The warning also appears as a banner in both reports. Either collect that node again first (example 2), or assess with the node list the evidence was collected with, plus the new entry:

```
git show 53081f6:nodes.yaml > /tmp/nodes-published.yaml
# append the eossc-example entry from step 1 to /tmp/nodes-published.yaml
uv run basic-check assess --nodes /tmp/nodes-published.yaml --only eossc-example \
--skip Italy --run live-2026-10-01-plus-eossc-example
```

Step 5: re-run the tests against the new published run.

```
uv run pytest -q
```

Two tests read the committed evidence. Each exists to make a change visible, not to be satisfied mechanically:

- `test_privacy.py::test_the_committed_evidence_publishes_no_personal_address_or_phone` must pass as it stands. If it fails, personal data is about to be published. Stop and find out why; never change the test to let it through.
- `test_names.py::test_the_official_names_match_the_nodes_that_show_them` pins which nodes display their approved name (four for the 26 September run; five on 24 September, when GÉANT's page was served). If the new page shows its name, add its id to the expected set and update the docstring to say why.
- `test_update_row.py::test_published_results_are_the_run_these_tests_copy` pins how many nodes the run has and which were skipped (13 and none, for run `web-5`). Update it to the new run's shape.

Step 6: update what quotes the published figures, then commit. `results/` is regenerated, but these are written by hand:

- a review document for the new run, like `results/REVIEW-2026-09-24.md`;
- the headline tally in `README.md` ("130 cells: 45 PASS · 8 FAIL · 77 review");
- "The published figures" in `docs/TEST-SUITE.md`, and the node list in "The node list — a YAML file" in the same file.

```
git status
git add nodes.yaml checklist/approved-names.txt checklist/approved-names-scoped.txt \
results README.md docs tests
git commit -m "Add eossc-example; publish run live-2026-10-01-plus-eossc-example"
git push
```

Example 2 — changing a node's landing page URL

The example is the change actually made on 25 September 2026: BBMRI-ERIC moved from <https://www.bbmri-eric.eu/eosc-node-bbmri-eric/> to <https://dev3.bbmri-eric.eu/eosc-node-bbmri-eric/>. It was set back on 26 September, because the `dev3.` host's `robots.txt` excludes the checker, so `nodes.yaml` has the `www.` address again.

Which files change. One line in one file: `url` for that node in `nodes.yaml`. Nothing else in the repository stores a node's URL. Add a comment recording the old address and the date, so the published evidence can still be traced:

```
- id: bbmri-eric
  name: BBMRI-ERIC
  # URL changed on 25 September 2026, from
  # https://www.bbmri-eric.eu/eosc-node-bbmri-eric/
  url: https://dev3.bbmri-eric.eu/eosc-node-bbmri-eric/
  eosc_page: https://eosc.eu/building-the-eosc-federation/eosc-node-bbmri-eric/
```

Change `eosc_page` as well only if the node's own entry on `eosc.eu` moved; a new landing page does not usually mean a new `eosc.eu` entry. The name files change only if the node's approved name changed. Keep the `id`: it names the evidence file, and changing it would make the node look new.

Step 1: check the configuration.

```
uv run pytest -q
```

Step 2: trial run for that node only. This contacts one site.

```
uv run basic-check collect --only bbmri-eric --results /tmp/trial
uv run basic-check assess --only bbmri-eric --results /tmp/trial
uv run basic-check show bbmri-eric --results /tmp/trial
```

Check that the page actually loaded before going further. On 25 September the `dev3.` host's `robots.txt` disallowed every path, the tool honoured it, and all ten points were `ERROR`. A result like that is not worth publishing: ask the node for its public address, or for the checker to be allowed, and stop there.

Step 3: update that node's row in the published run. Only BBMRI-ERIC is fetched again; the other rows are copied from `results/results.json` as they are, neither re-fetched nor re-judged:

```
uv run basic-check run --update-results-for-node bbmri-eric
```

That fetches the one page into a scratch directory, and only if the page loaded does it replace `bbmri-eric`'s evidence, screenshot and row, then rebuild `results.md`, `index.html` and `results.csv` from the stored rows. A failed fetch changes nothing (add `--accept-error` to record it anyway). It prints each point's old and new verdict, and both reports gain an **Updated rows** banner naming the node, the date and the previous URL. The run id and time in the header stay those of the original run. To reuse the evidence from step 2 instead of fetching again, copy it into `results/evidence/` and run `uv run basic-check assess --update-results-for-node bbmri-eric`, which is offline. To see the outcome before touching the published run, try it on a copy first: `cp -r results /tmp/copy` and add `--results /tmp/copy`. For a completely fresh table instead, follow path B of example 1 step 4.

Why the order matters. `assess` labels each row with the URL in `nodes.yaml`, but takes the verdicts from the evidence on disk. Between step 1 and step 3, the two describe different pages. `assess` now detects this. It prints a warning naming the node and both URLs, puts an **Evidence from a different URL** banner in `results.md`

and `index.html`, and records the pair under `url_mismatch` in `results.json`. Step 3 clears it, because the evidence now comes from the configured URL.

To rebuild a run collected with an older node list, without the warning, use that list. For the reviewed run of 24 September (in the git history at `800d632`), assessed in a scratch copy of its `results/`:

```
git show 53081f6:nodes.yaml > /tmp/nodes-2026-09-24.yaml
uv run basic-check assess --run live-2026-09-24-no-italy --skip Italy \
--nodes /tmp/nodes-2026-09-24.yaml
```

Step 4: tests, review, commit. As in example 1, steps 5 and 6: run `uv run pytest -q`, check the two tests that read the committed evidence, write the review, update the quoted figures, and commit `nodes.yaml`, `results/` and the documents together.

Example 3 — moving to a new revision of the checklist

There is no URL to change. The tool never downloads the checklist document. A copy of it is committed in `checklist/`, and three things connect the tool to it:

Where	What it says
<code>src/basic_check/cli.py</code> , the line <code>DEFAULT_CHECKLIST = ROOT / "checklist" / "v3.1.yaml"</code>	Which checklist revision is used when <code>--checklist</code> is not given. This is the one line to change to make a new revision the default. The help texts, the tests and the reports all follow it.
<code>checklist/vX.Y.yaml</code> : <code>checklist_version</code> , <code>checklist_date</code> , <code>source_document</code> , <code>source_file</code> , <code>source_sha256</code>	Which document that revision was transcribed from, and its SHA-256 hash.
<code>--checklist / -c <file></code> on <code>assess</code> , <code>run</code> and <code>points</code>	A different revision for a single run, without changing the default.

When the originating document changes, you therefore add a new file for the new revision and point the default at it. `v3.1.yaml` is never edited: the published reports name the file they were produced from.

The example assumes a v3.2 dated 15 October 2026. Use the real version, date and filenames. This procedure was last applied on 26 September 2026, to move from v3.0 to v3.1; the header of `checklist/v3.1.yaml` records what changed, point by point.

Step 1: commit the new document next to the old one.

```
cp ~/Downloads/20261015_Node_Landing_Page_Verification_Checklist_v3.2.pdf checklist/
sha256sum checklist/20261015_Node_Landing_Page_Verification_Checklist_v3.2.pdf
# on macOS: shasum -a 256 checklist/20261015_Node_Landing_Page_Verification_Checklist_v3.2.pdf
```

The committed v3.1 file is a PDF rendering of the circulated `.docx`. Commit the `.docx` itself if you prefer: `source_file` may name either.

Step 2: create the new revision's file from the old one.

```
cp checklist/v3.1.yaml checklist/v3.2.yaml
```

In `checklist/v3.2.yaml`, change the header:

```
checklist_version: "3.2"
checklist_date: "2026-10-15"
```

```
source_document: 20261015_Node_Landing_Page_Verification_Checklist_v3.2.docx
source_file: 20261015_Node_Landing_Page_Verification_Checklist_v3.2.pdf
source_sha256: <the hash printed in step 1>
```

The version must match the filename (`v3.2.yaml` declares "3.2"); a test enforces this.

Step 3: compare every point with the new document. No tool can do this step for you. For each point in `v3.2.yaml`:

- If the wording changed, update `title` and `requirement`, which are quoted word for word. Then read the `check_*` function named in `implemented_by`, in `src/basic_check/checks.py`, and change it if it no longer applies the new wording.
- If what a script can decide changed, update `decidable` and `decidable_note`.
- If a point was **added**, write a new `check_<id>` function in `checks.py`, call it from `run_all()` in `checklist` order, name it in the point's `implemented_by`, and add tests for it to `tests/test_checks.py`.
- If a point was **removed**, delete its entry, its function and its tests.

Step 4: try the new revision offline, before switching. `-c` applies it to one run only. Re-scoring saved evidence contacts no one:

```
uv run basic-check points -c checklist/v3.2.yaml
mkdir -p /tmp/v32 && cp -r results/evidence /tmp/v32/
uv run basic-check assess -c checklist/v3.2.yaml --results /tmp/v32 --run trial-v3.2
```

Compare `/tmp/v32/results.md` with `results/results.md`. Any difference comes from the new rules alone, since the evidence is the same. If the published run skipped a node or used an older `nodes.yaml`, pass the same `--skip` and `--nodes` here.

Step 5: make it the default, and run the tests. In `src/basic_check/cli.py`:

```
DEFAULT_CHECKLIST = ROOT / "checklist" / "v3.2.yaml"
uv run pytest -q
uv run basic-check --help # now says "against checklist v3.2"
```

The tests check the new file's hash, that its version matches its filename, and that every point maps to exactly one `check_*` function and every function to one point. They also check that `v3.1.yaml` and its document are still intact, so older runs can still be rebuilt with `-c checklist/v3.1.yaml`. What they cannot check is whether a rule is right: that is step 3.

Step 6: produce and publish the results. The default revision is now `v3.2`, so no `-c` is needed. Choose between the same two paths as in example 1:

```
# fresh evidence from every node:
uv run basic-check run --results /tmp/new-run --run live-2026-10-20-v3.2
# review, then:
rm -rf results/evidence && cp -r /tmp/new-run/evidence results/evidence
uv run basic-check assess --run live-2026-10-20-v3.2
```

Alternatively, re-score the published evidence against `v3.2` without contacting anyone, as in step 4 but writing to `results/`. The report then says `v3.2`, while the capture date is still that of the evidence.

The reports now link to `results/checklist-v3.2.html`. The old `results/checklist-v3.1.html` stays behind; remove it, and update the places that name `v3.1` in prose or links:

```
git rm results/checklist-v3.1.html
rg -n "v3\.1" README.md docs checklist/README.md .github
```

In `README.md`, the table of checklist points links each point to `results/checklist-v3.1.html#p....`. Those links break once the file is removed, so change them to `v3.2`.

Step 7: tests, review, commit. As in example 1, steps 5 and 6. If any `check_*` function changed in step 3, also update its decision procedure in section 5 of `ANALYSIS-WORKFLOW.md`, which describes each point's rule branch by branch. Commit the new document, `checklist/v3.2.yaml`, `cli.py`, any changed checks and tests, `results/` and the documents together, so the commit shows the whole change.

Sources

- Checklist: [checklist/20260910_Node_Landing_Page_Verification_Checklist_v3.1.pdf](#), transcribed to [checklist/v3.1.yaml](#) with its SHA-256 pinned
- EOSC Federation node index: <https://eosc.eu/building-the-eosc-federation/>
- Repository: <https://github.com/marioarele/eosc-basic-compliance>
- `uv` documentation: <https://docs.astral.sh/uv/>
- Playwright for Python: <https://playwright.dev/python/>